

Discussion 8

B+Tree & Sorting

EECS 484

Logistics

- **Project 3** due **Today** at 11:45 PM ET
- **HW 4** due **Nov 8th** at 11:45 PM ET
- **HW 5** released, due **Nov 15th** at 11:45 PM ET

B+Trees

B+ Trees

- Self balancing tree structure with multiple elements in each node
 - All leaf nodes are the same height/depth
 - Height = length of any path from root to the leaf
 - A B+Tree is an M-way search tree
 - Every node other than the root is at least half-full $M/2 - 1 \leq \#keys \leq M - 1$
 - Every inner node with k keys has k+1 non-null children
 - Max fanout = M
 - Max pointers in an inner node (maximum number of children for a node)
- 3 main operations
 - Search
 - Insert
 - Delete
 - B+Tree Use: Increase speed of lookups based on an attribute(s) in your table to improve efficiency of these operations in a large DB

B+ Tree Leaf Node Values

Keys: Based on attribute(s) that the index is based on

Values:

Approach #1: Record IDs


- A pointer to the location of the tuple to which the index entry corresponds.

Approach #2: Tuple Data

- The leaf nodes (of the primary key index) store the actual contents of the tuple.
- Secondary indexes must store the Record ID as their values.


Searching in B+ Tree

- Searching for a particular element
 - The DBMS can use a B+Tree index if the query provides any of the attributes of the search key.
 - Follow the pointers in each node until you find the leaf the element SHOULD exist in
 - No guarantee, if it doesn't exist in the leaf node it doesn't exist in the tree
 - Pointers are “guides”
 - “If you're looking for less than 17, this way, else that way”




Inserting into a B+ Tree

- Add an element to the correct leaf node
 - If the desired leaf node has capacity, easy
 - Otherwise need to either split or redistribute
- Normal Insert
 - Insert 71*




Inserting into a B+ Tree

- Add an element to the correct leaf node
 - If the desired leaf node has capacity, easy
 - Otherwise need to either split or redistribute
- Normal Insert
 - Insert 71^*




Inserting into a B+ Tree

- Add an element to the correct leaf node
 - If the desired leaf node has capacity, easy
 - Otherwise need to either split or redistribute
- Normal Insert
 - Insert 71^*




Inserting into a B+ Tree

- Redistribute elements to left sibling
 - Insert 20*




Inserting into a B+ Tree

- Redistribute elements to left sibling
 - Insert 20*



Inserting into a B+ Tree


- Redistribute elements to left sibling
 - Insert 20*



Kick smallest element left

Inserting into a B+ Tree

- Redistribute elements to left sibling
 - Insert 20*




Kick smallest element left

Inserting into a B+ Tree


- Redistribute elements to left sibling
 - Insert 20*

No longer correct




Inserting into a B+ Tree

- Redistribute elements to left sibling
 - Insert 20*




Inserting into a B+ Tree

- Redistribute elements to left sibling
 - Insert 20*
 - And we're done :)




Inserting into a B+ Tree

- Split with extra elements in right child
 - Insert 20*




Inserting into a B+ Tree

- Split with extra elements in right child
 - Insert 20*




Inserting into a B+ Tree

- Split with extra elements in right child
 - Insert 20*




Inserting into a B+ Tree

- Split with extra elements in right child
 - Insert 20*



Inserting into a B+ Tree

- Split with extra elements in right child
 - Insert 20*




Promote middle key and split

Inserting into a B+ Tree

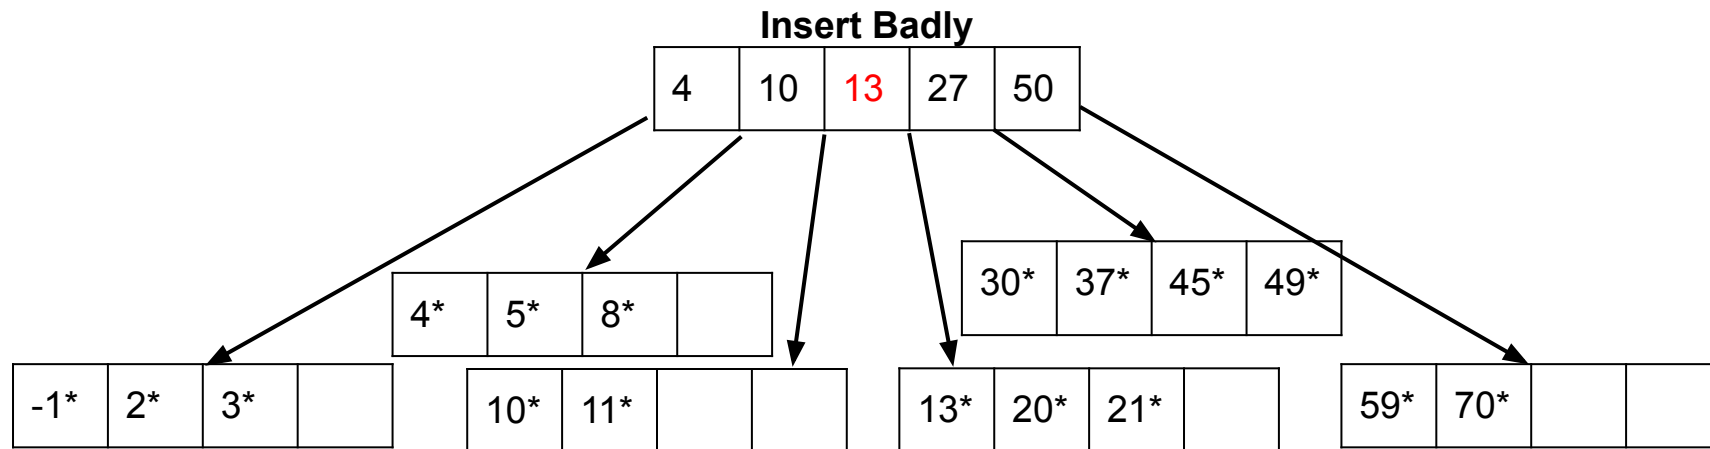
- Split with extra elements in right child
 - Insert 20*

Full :(



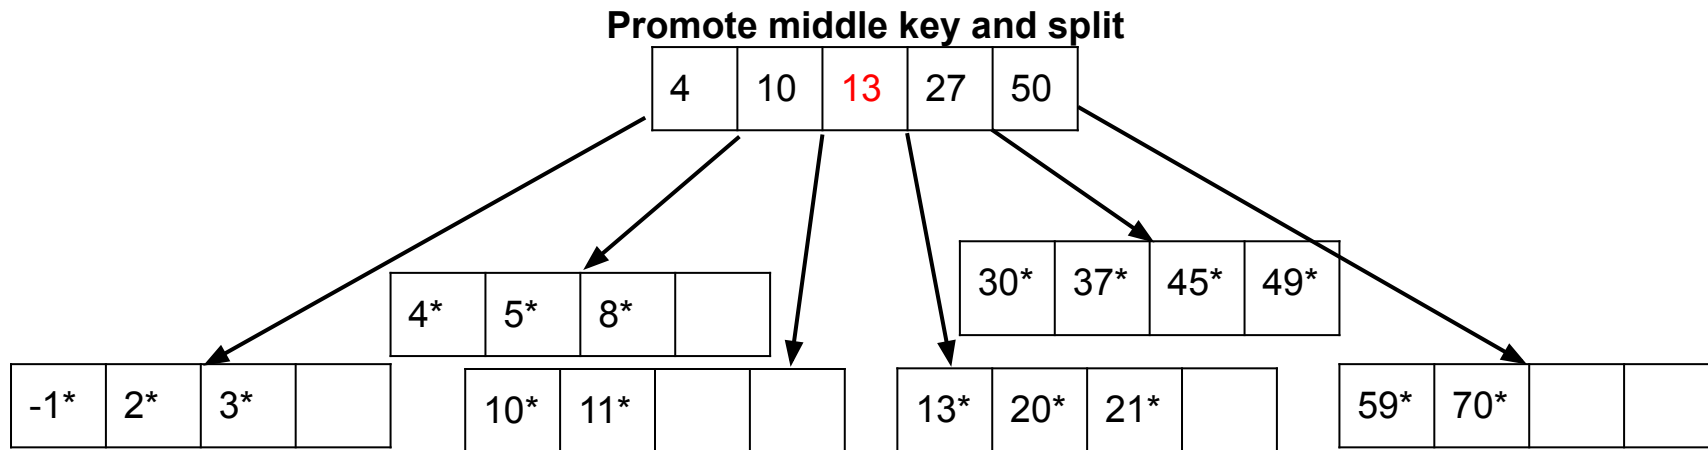
Inserting into a B+ Tree

- Split with extra elements in right child
 - Insert 20*




Inserting into a B+ Tree

- Split with extra elements in right child
 - Insert 20*



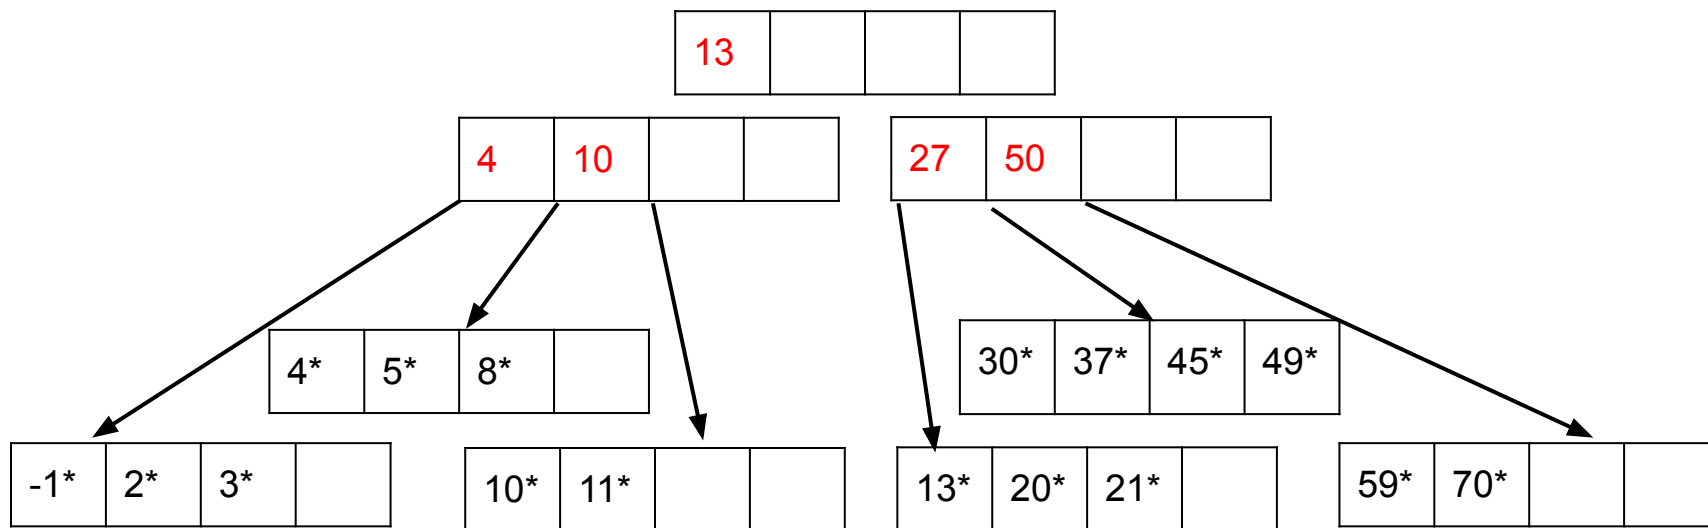
Inserting into a B+ Tree

- Split with extra elements in right child
 - Insert 20*



Inserting into a B+ Tree


- Split with extra elements in right child
 - Insert 20*



Inserting into a B+ Tree


- Split with extra elements in right child
 - Insert 20*

Reconnect



Inserting into a B+ Tree

- Split with extra elements in right child
 - Insert 20*
 - Finally done!




Inserting into a B+ Tree

Takeaways

- Redistributing is a lot less work
 - Usually smaller height
 - More data entries per page
 - More I/O (need to check right/left nodes)
 - Can't do this if the right and left nodes are full


Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Normal Delete
 - Delete 21*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Normal Delete
 - Delete 21*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Normal Delete
 - Delete 21*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Normal Delete
 - Delete 21*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Normal Delete
 - Delete 21*
 - Done




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*



Deleting from a B+ Tree


- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*



Num elements < $M/2-1$

Deleting from a B+ Tree


- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*



Borrow from left


Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try redistribution
 - Delete 59*
 - Done :)



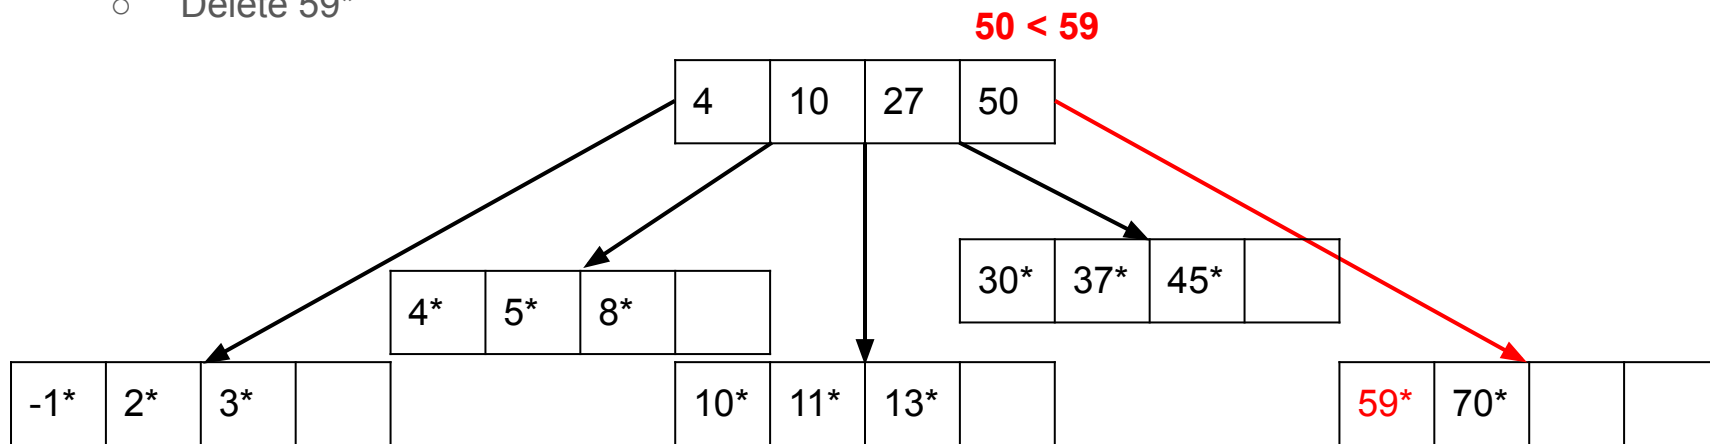
Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try merging
 - Delete 59*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try merging
 - Delete 59*




Deleting from a B+ Tree

- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try merging
 - Delete 59*



Deleting from a B+ Tree


- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try merging
 - Delete 59*



Num elements < $M/2-1$

Deleting from a B+ Tree


- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try merging
 - Delete 59*



Merge with left

Deleting from a B+ Tree


- Delete an element from the tree
 - If the leaf node is at least half-full, then easy
 - Otherwise need to either redistribute or merge
- Try merging
 - Delete 59*
 - Done :)



Clustered vs Unclustered Index

A clustered index means the table is stored in the sort order specified by the primary key.

Retrieving tuples in the order they appear in a unclustered index can be very inefficient.



Index Matching

- When is it appropriate to use an index to evaluate a selection predicate?
 - Conjunction (ANDs) of terms involving only attributes (no disjunctions, ORs)
 - Hash Index
 - Only equality operation, predicate has all index attributes
 - Tree index
 - Any operations, attributes are a prefix of the search key
 - Works in other cases as well, but may be less common and less efficient

Index Matching

Considering an index on $\langle a, b, c \rangle$ in the following situations, will a tree index be more efficient than a file scan? Is it appropriate to use a hash index?

Selection Predicate	Tree Index	Hash Index
$a=5$ and $b=3$		
$a>5$ and $b<3$		
$b=3$		
$a=7$ and $b=5$ and $c=4$ and $d>4$		
$a=7$ and $c=5$		

Index Matching

Considering an index on $\langle a, b, c \rangle$ in the following situations, will a tree index be more efficient than a file scan? Is it appropriate to use a hash index?

Selection Predicate	Tree Index	Hash Index
$a=5$ and $b=3$	Yes	No
$a>5$ and $b<3$		
$b=3$		
$a=7$ and $b=5$ and $c=4$ and $d>4$		
$a=7$ and $c=5$		

Index Matching

Considering an index on $\langle a, b, c \rangle$ in the following situations, will a tree index be more efficient than a file scan? Is it appropriate to use a hash index?

Selection Predicate	Tree Index	Hash Index
$a=5$ and $b=3$	Yes	No
$a>5$ and $b<3$	Yes	No
$b=3$		
$a=7$ and $b=5$ and $c=4$ and $d>4$		
$a=7$ and $c=5$		

Index Matching

Considering an index on $\langle a, b, c \rangle$ in the following situations, will a tree index be more efficient than a file scan? Is it appropriate to use a hash index?

Selection Predicate	Tree Index	Hash Index
$a=5$ and $b=3$	Yes	No
$a>5$ and $b<3$	Yes	No
$b=3$	No	No
$a=7$ and $b=5$ and $c=4$ and $d>4$		
$a=7$ and $c=5$		

Index Matching

Considering an index on $\langle a, b, c \rangle$ in the following situations, will a tree index be more efficient than a file scan? Is it appropriate to use a hash index?

Selection Predicate	Tree Index	Hash Index
$a=5$ and $b=3$	Yes	No
$a>5$ and $b<3$	Yes	No
$b=3$	No	No
$a=7$ and $b=5$ and $c=4$ and $d>4$	Yes	Yes
$a=7$ and $c=5$		

Index Matching

Considering an index on $\langle a, b, c \rangle$ in the following situations, will a tree index be more efficient than a file scan? Is it appropriate to use a hash index?

Selection Predicate	Tree Index	Hash Index
$a=5$ and $b=3$	Yes	No
$a>5$ and $b<3$	Yes	No
$b=3$	No	No
$a=7$ and $b=5$ and $c=4$ and $d>4$	Yes	Yes
$a=7$ and $c=5$	Yes	No

External Sorting

External Sorting

- Sorting is nice
 - We have lots of nice algorithms that will sort for us
 - Quicksort, Mergesort, Heapsort, etc.
 - We can do this very quickly with lots of data - $O(N \log N)$
- But what if we have too much data to fit in RAM?
 - We can still sort but it will be so so slow :(
 - Need some way to *externally* sort the data on the disk while dealing with limited fast memory


General External Merge Sort

- Step 1:
 - Have a large dataset of N pages that you would like to sort using B buffer pages
- Step 2:
 - Divide the dataset into $\lceil N/B \rceil$ runs (each of which is B pages long)
- Step 3:
 - Sort each run by itself normally using your favorite algorithm
 - We can fit the entire run of B pages into our RAM so no problem
- Step 4:
 - Sort the runs amongst each other
 - We can merge $B-1$ runs at a time
 - $B-1$ pages for each run plus 1 page to store the output
 - Each run is larger than 1 page though!
 - Load the first (sorted) page of each run and once it's empty, read the next page
 - Similarly, write the output buffer each time we run out of space and keep going

Step 1

- Have a dataset


Suppose $B = 4$ and
each page can hold
2 bars in full.



Step 2


- Divide the data into $\text{ceiling}(N/B)$ runs
 - Each is B pages long, i.e. each run is technically supposed to have 8 bars
 - (for simplicity we only show 4 smallest bars in each run)

Suppose $B = 4$ and each page can hold 2 bars in full.




Step 3

- Sort each run individually (for simplicity we only show 4 smallest bars in each run)



Step 4

- Sort the runs with each other
 - B-1 runs at a time




(for simplicity we only show 4 smallest bars in each run)


Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - B-1 runs at a time



Load 1 (sorted) page at
a time from each run



Single output page




(for simplicity we only show 4
smallest bars in each run)

Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)


Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)



Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)


Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)



Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)



Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




Output page full



(for simplicity we only show 4
smallest bars in each run)


Step 4

- Sort the runs with each other
 - B-1 runs at a time



(for simplicity we only show 4 smallest bars in each run)


Take minimum element
from all loaded pages
Remember Merge Sort



Write to disk
Empty page and continue




Suppose $B = 4$ and
each page can hold
2 bars in full.





Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.


- Sort the runs with each other
 - $B-1$ runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)





Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.


- Sort the runs with each other
 - B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort



(for simplicity we only show 4
smallest bars in each run)




Step 4




Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other


- B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)




Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.


- Sort the runs with each other
 - $B-1$ runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)





Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.


- Sort the runs with each other
 - B-1 runs at a time



Continue till runs sorted
Will need to make more
passes if more than B-1 runs



(for simplicity we only show 4
smallest bars in each run)



General External Merge Sort Math

- We have a dataset with N pages
 - We'll use B buffer pages
 - We'll have $\lceil N/B \rceil$ runs initially
 - We need to make passes over the runs until entire dataset is sorted
 - We merge $B-1$ runs together at a time
 - That means we have $\lceil \lceil N/B \rceil / (B-1) \rceil$ merged runs afterwards
 - Each time we make a pass we've merged all runs in sets of size $B-1$
 - We must continue to do this till we have 1 output dataset in sorted order
 - Takes $1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil$ passes
 - Total IO cost is $\# \text{passes} * 2N$
 - Each pass we read each page and write each page in a new sorted order

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(\text{ceiling}(N/B))) = 1 + \text{ceiling}(\log_{17}(50))$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(\text{ceiling}(N/B))) = 1 + \text{ceiling}(\log_{17}(50)) = 1 + \text{ceiling}(1.38\text{ish}) = 1 + 2 = 3$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(N/B)) = 1 + \text{ceiling}(\log_{17}(50)) = 1 + \text{ceiling}(1.38\text{ish}) = 1 + 2 = 3$
 - How many IO operations?

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(N/B)) = 1 + \text{ceiling}(\log_{17}(50)) = 1 + \text{ceiling}(1.38\text{ish}) = 1 + 2 = 3$
 - How many IO operations?
 - $\text{\#IO} = 2N * \text{\#passes}$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(N/B)) = 1 + \text{ceiling}(\log_{17}(50)) = 1 + \text{ceiling}(1.38\text{ish}) = 1 + 2 = 3$
 - How many IO operations?
 - $\text{\#IO} = 2N * \text{\#passes} = 2 * 900 * 3 = 5400$

Get started on Homework 4!

We're here if you need any help!!

- Office Hours: Schedule is [here](#), both virtual and in person offered
- Piazza
- Next week's discussion!!!