

University of Michigan

23 Two-Phase Locking



Database Management Systems
EECS 484
Fall 2024

LM

Lin Ma
Computer Science and
Engineering Division

LAST CLASS

Conflict Serializable

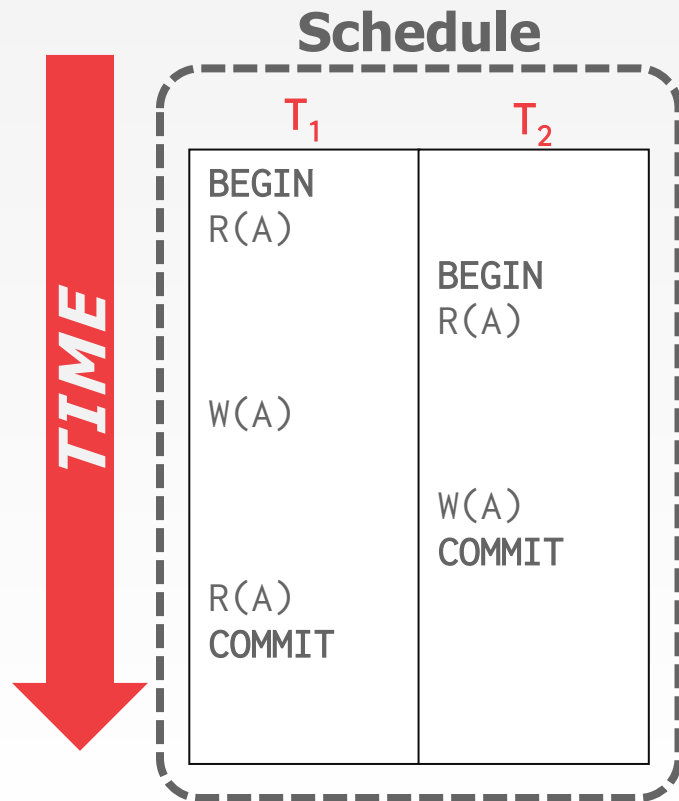
- Verify using either the "swapping" method or dependency graphs.
- Any DBMS that says that they support "serializable" isolation does this.

View Serializable

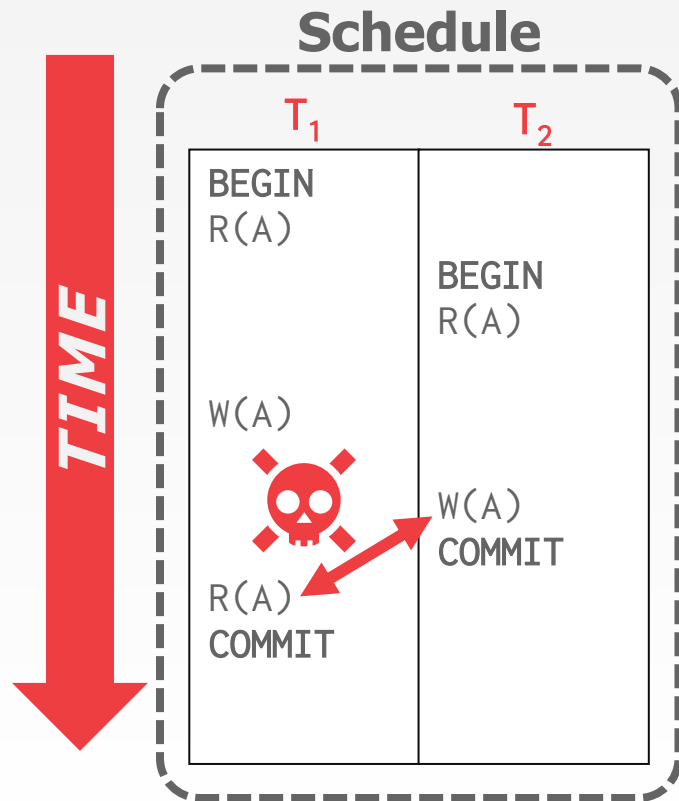
- No efficient way to verify.
- Lin doesn't know of any DBMS that supports this.



EXAMPLE



EXAMPLE



OBSERVATION

We need a way to guarantee that all execution schedules are correct (i.e., serializable) without knowing the entire schedule ahead of time.



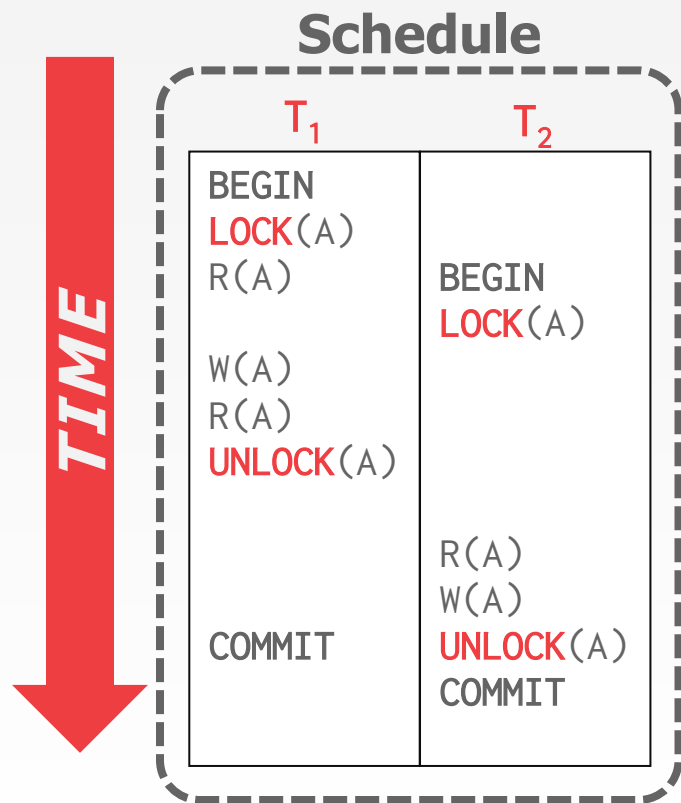
OBSERVATION

We need a way to guarantee that all execution schedules are correct (i.e., serializable) without knowing the entire schedule ahead of time.

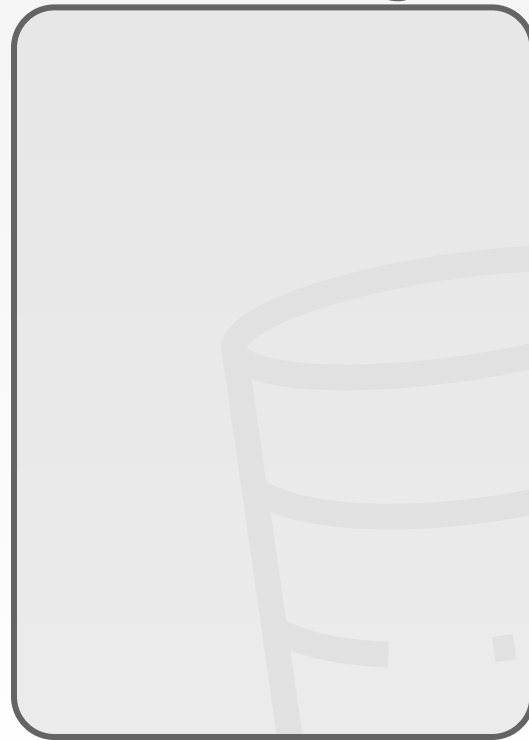
Solution: Use locks to protect database objects.



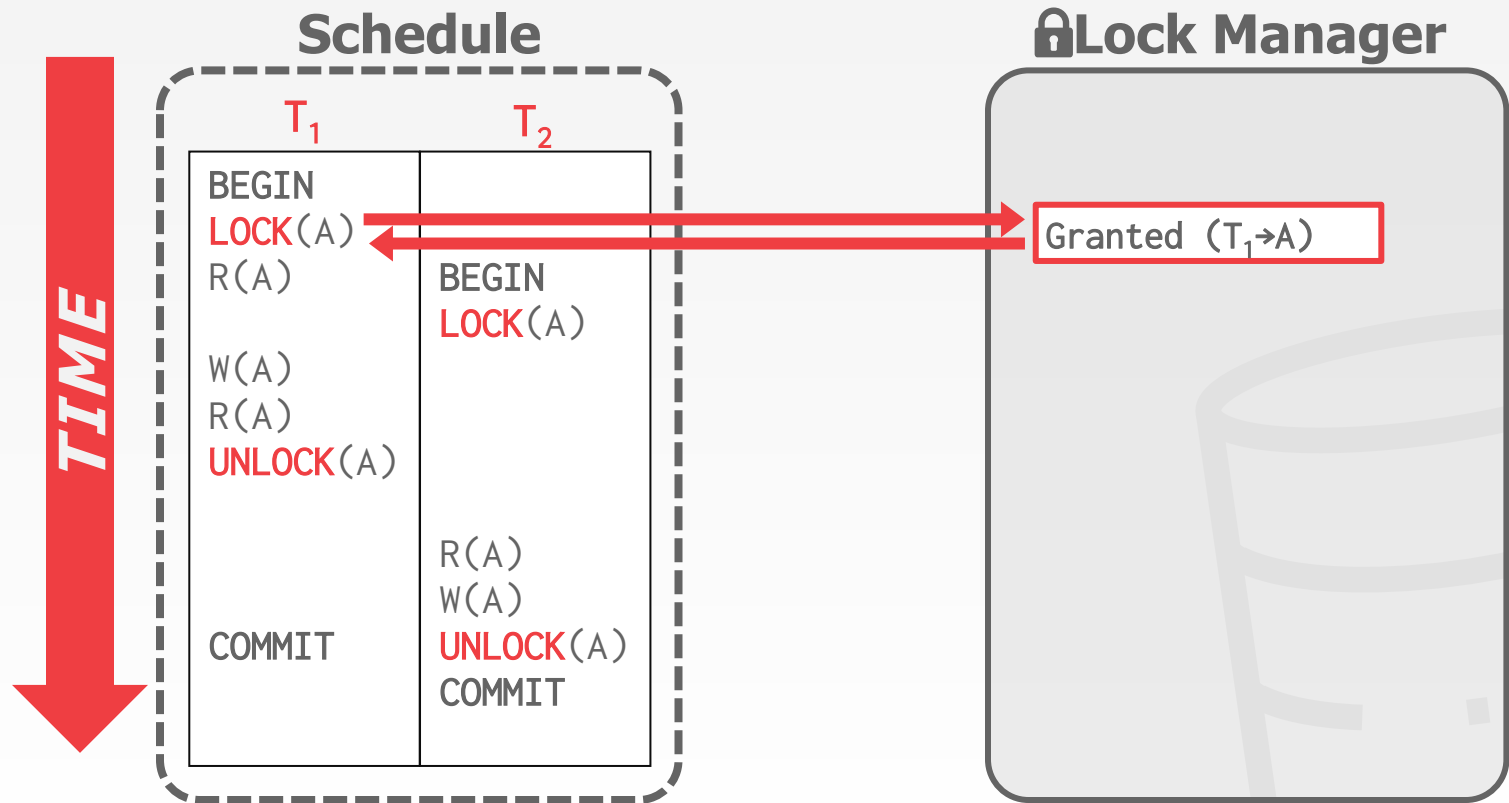
EXECUTING WITH LOCKS



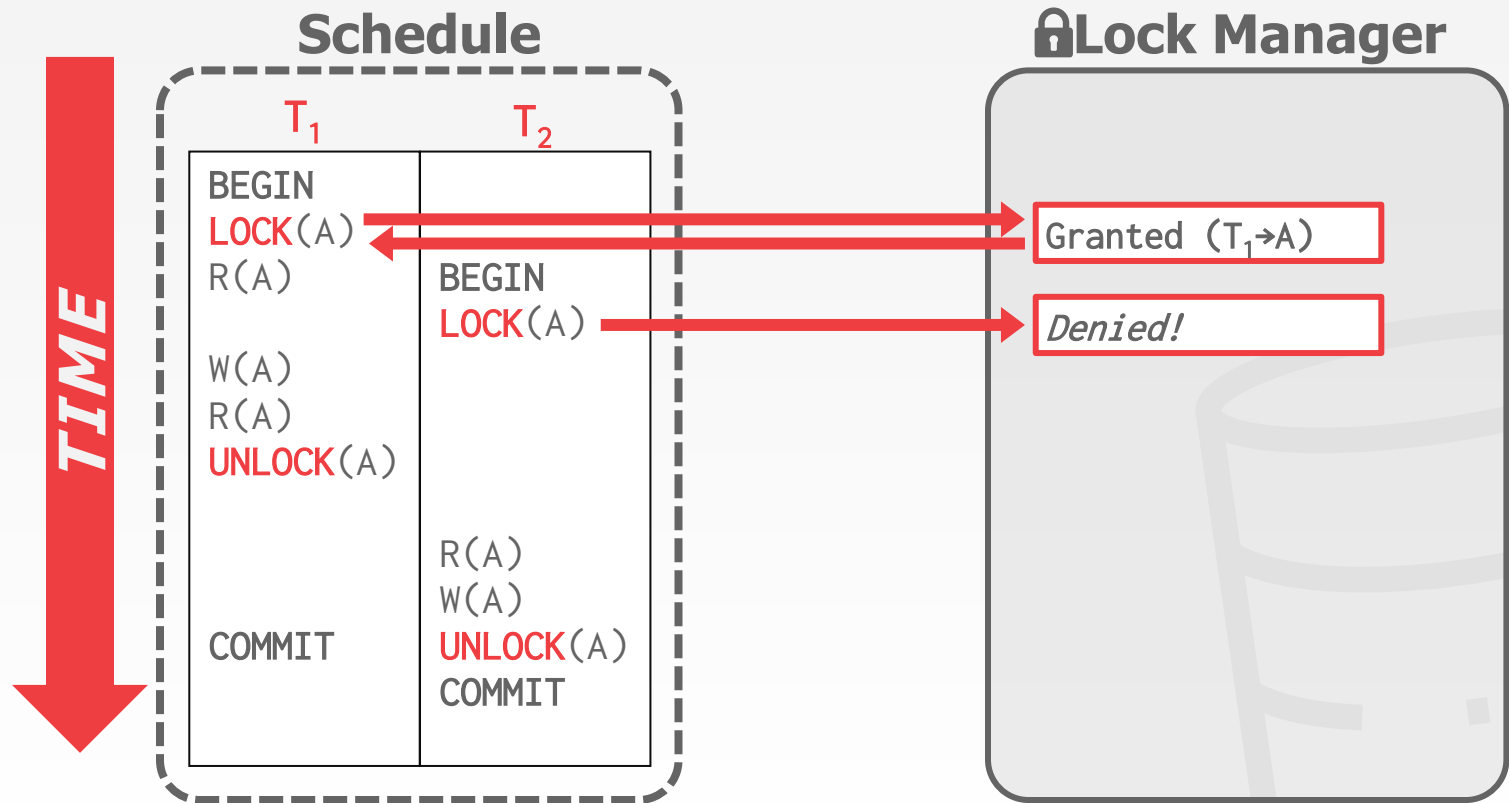
Lock Manager



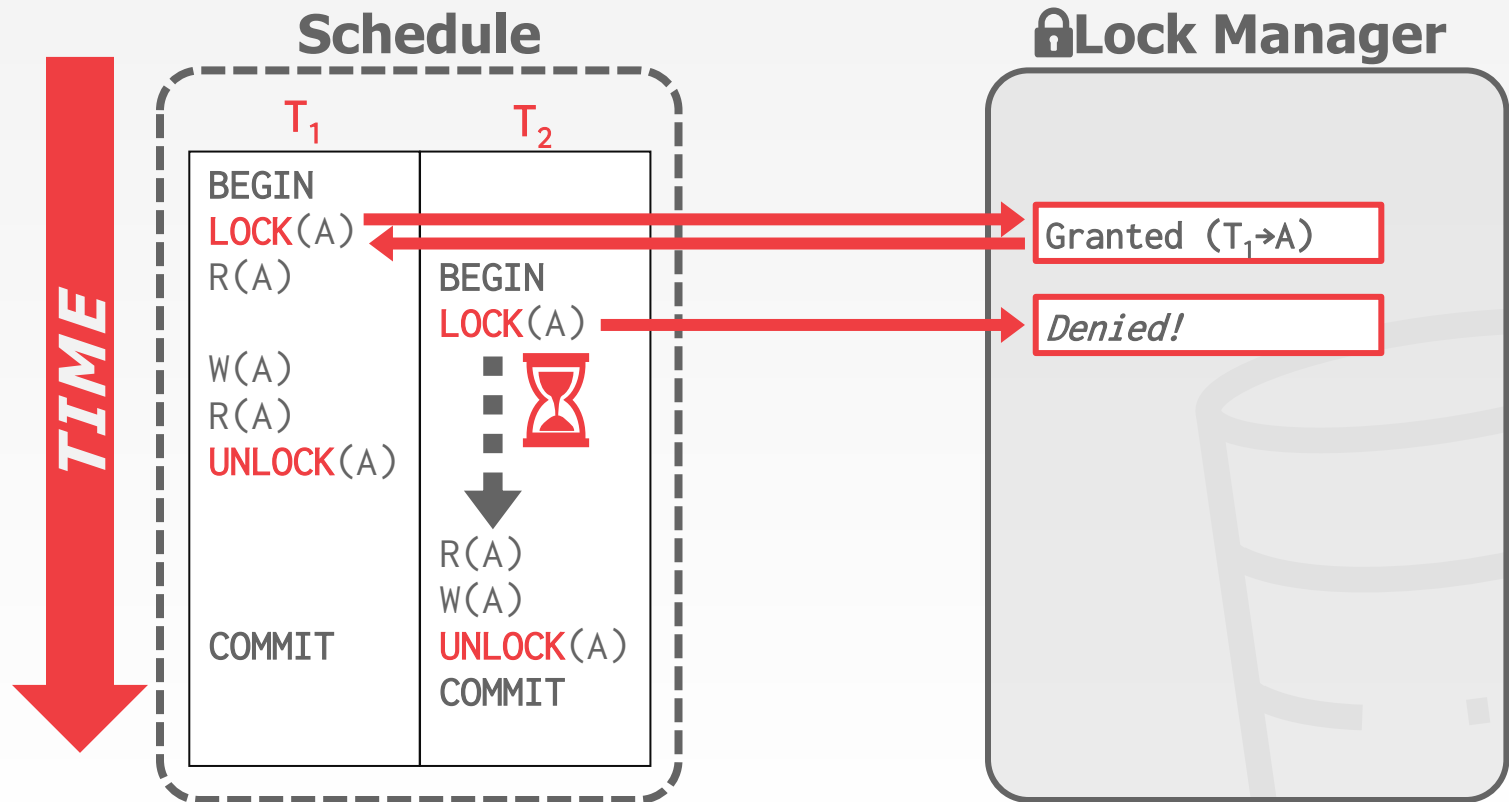
EXECUTING WITH LOCKS



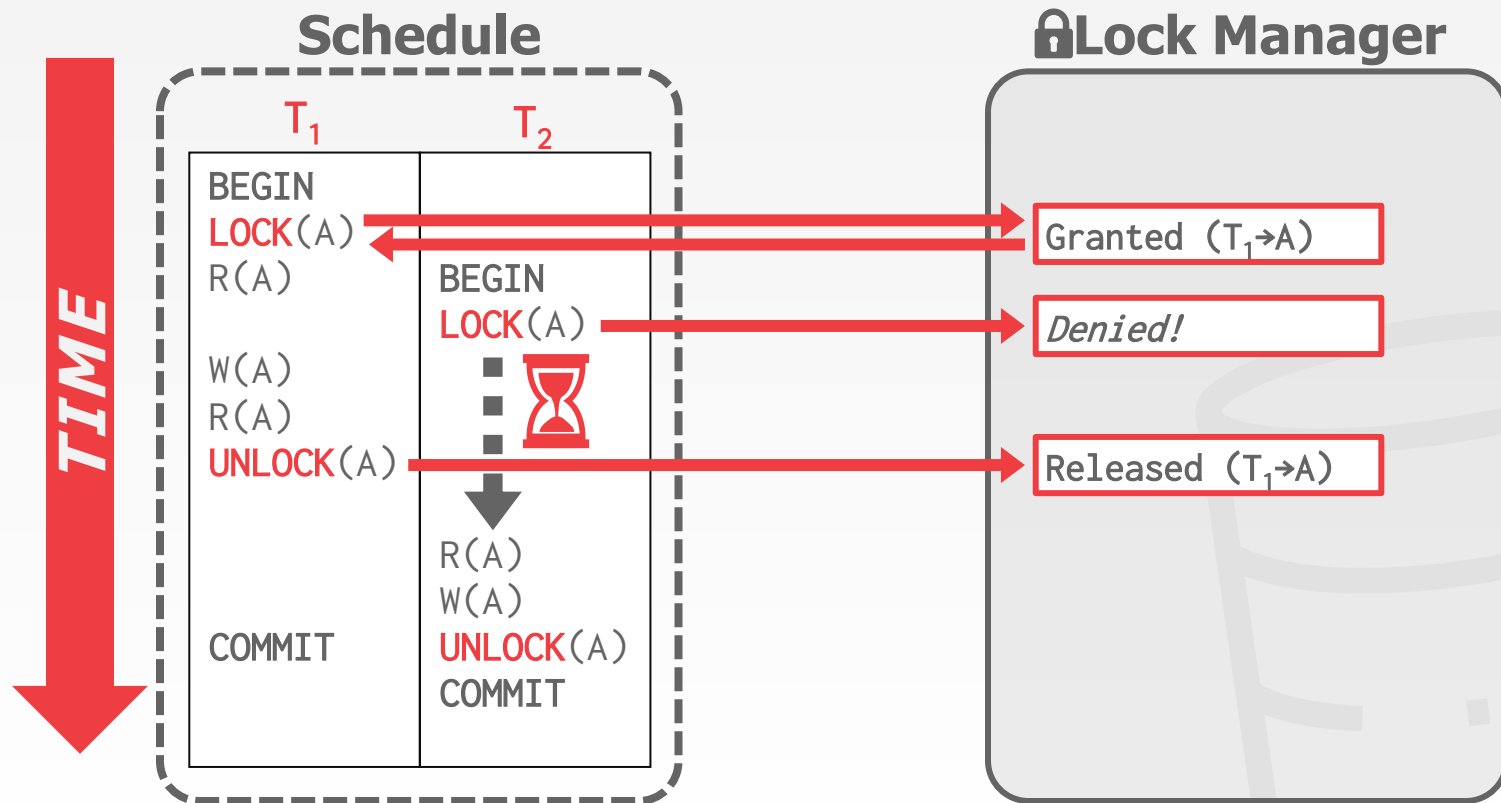
EXECUTING WITH LOCKS



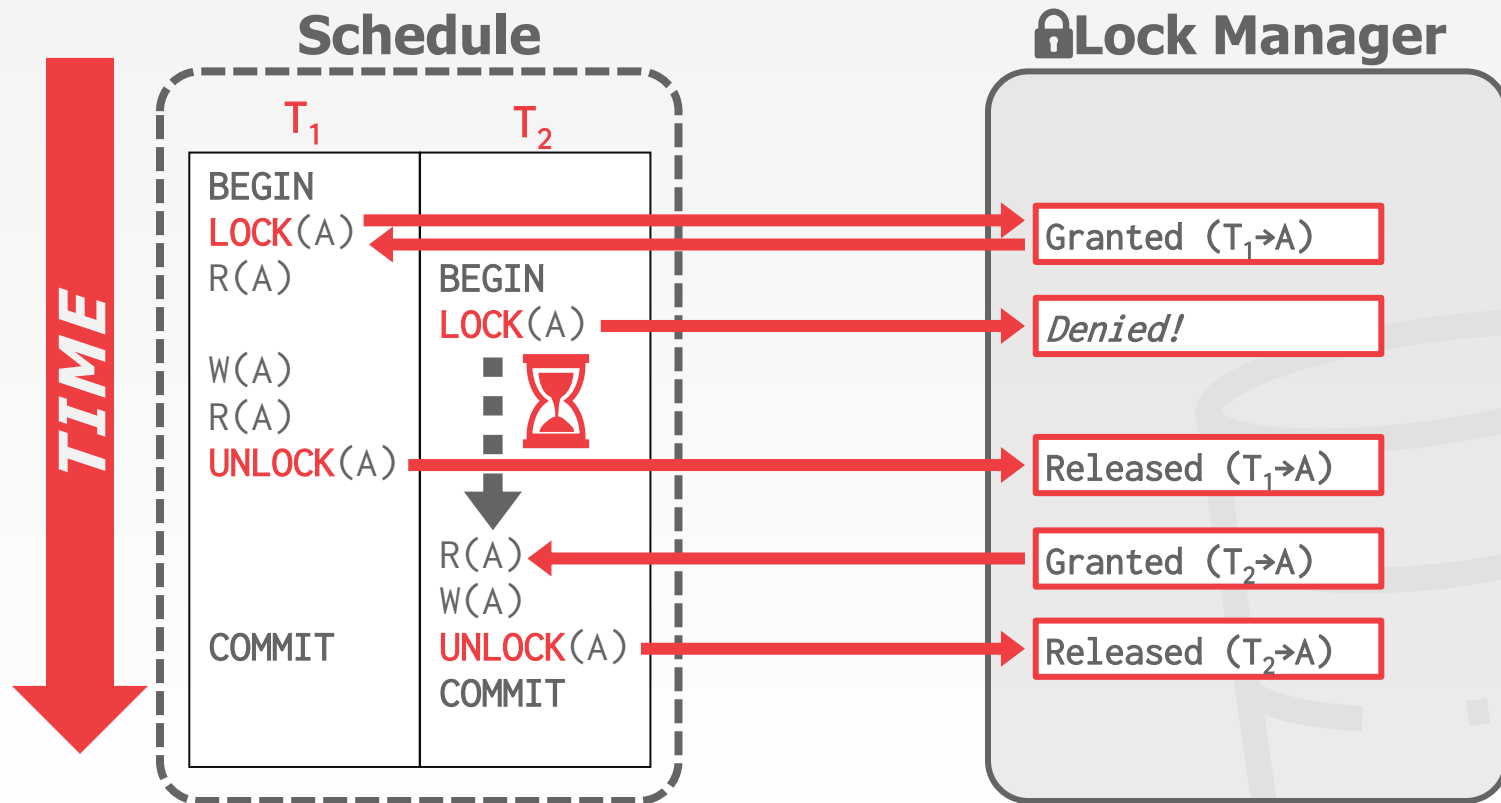
EXECUTING WITH LOCKS



EXECUTING WITH LOCKS



EXECUTING WITH LOCKS



TODAY'S AGENDA

Lock Types

Two-Phase Locking

Deadlock Detection + Prevention

Isolation Levels



BASIC LOCK TYPES

S-LOCK: Shared locks for reads.

X-LOCK: Exclusive locks for writes.

Compatibility Matrix		
	Shared	Exclusive
Shared	✓	X
Exclusive	X	X



EXECUTING WITH LOCKS

Transactions request locks (or upgrades).

Lock manager grants or blocks requests.

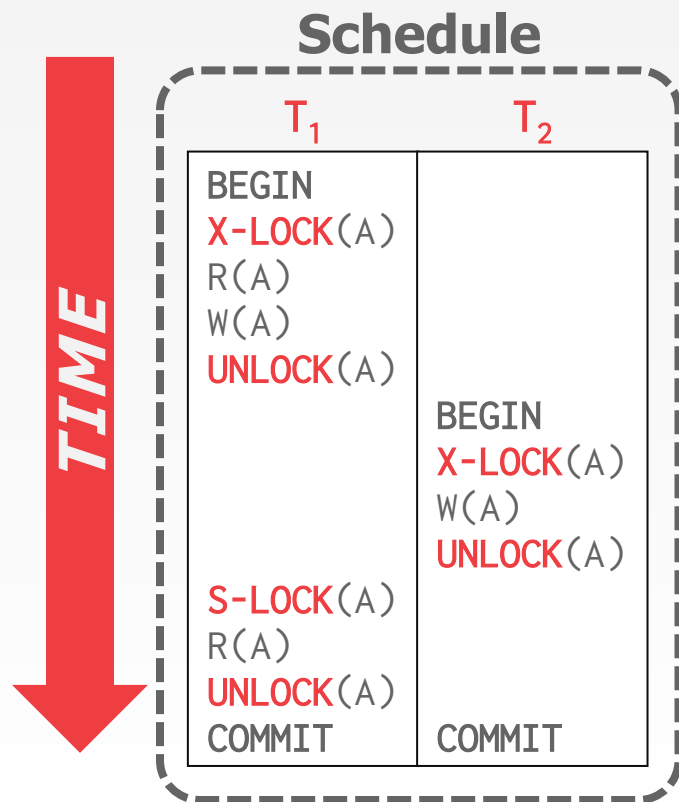
Transactions release locks.

Lock manager updates its internal lock-table.

→ It keeps track of what transactions hold what locks and what transactions are waiting to acquire any locks.

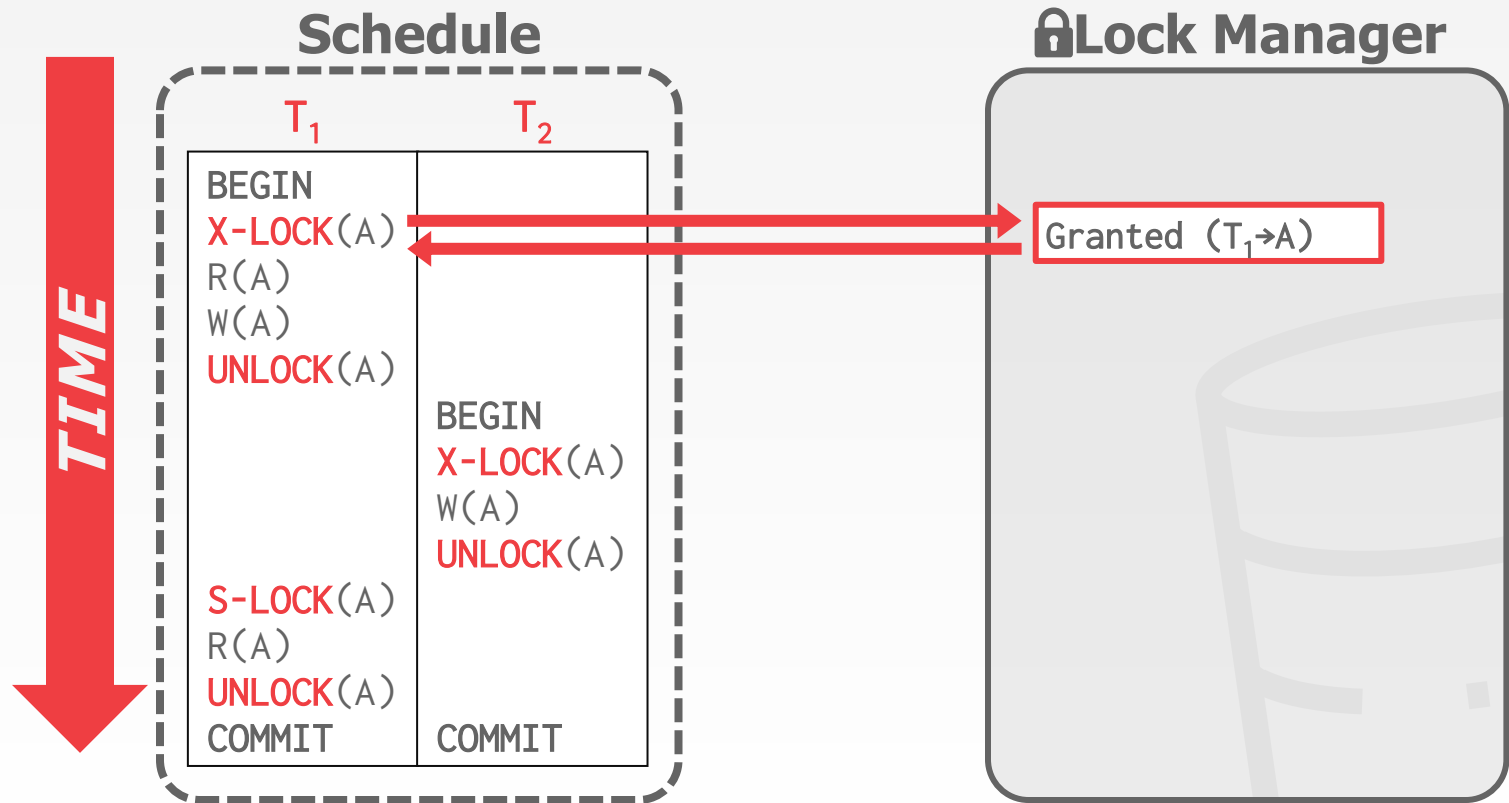


EXECUTING WITH LOCKS

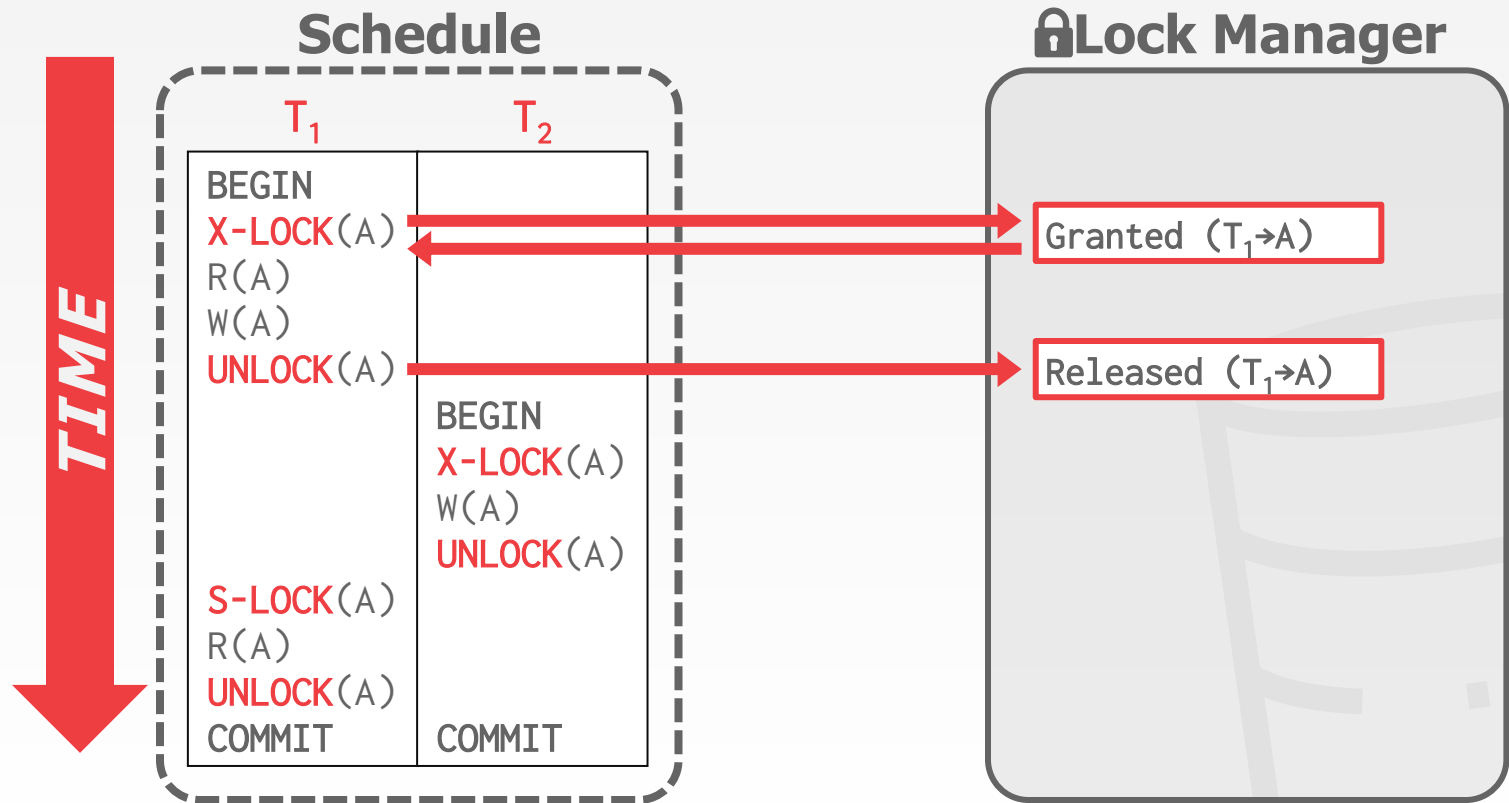


Lock Manager

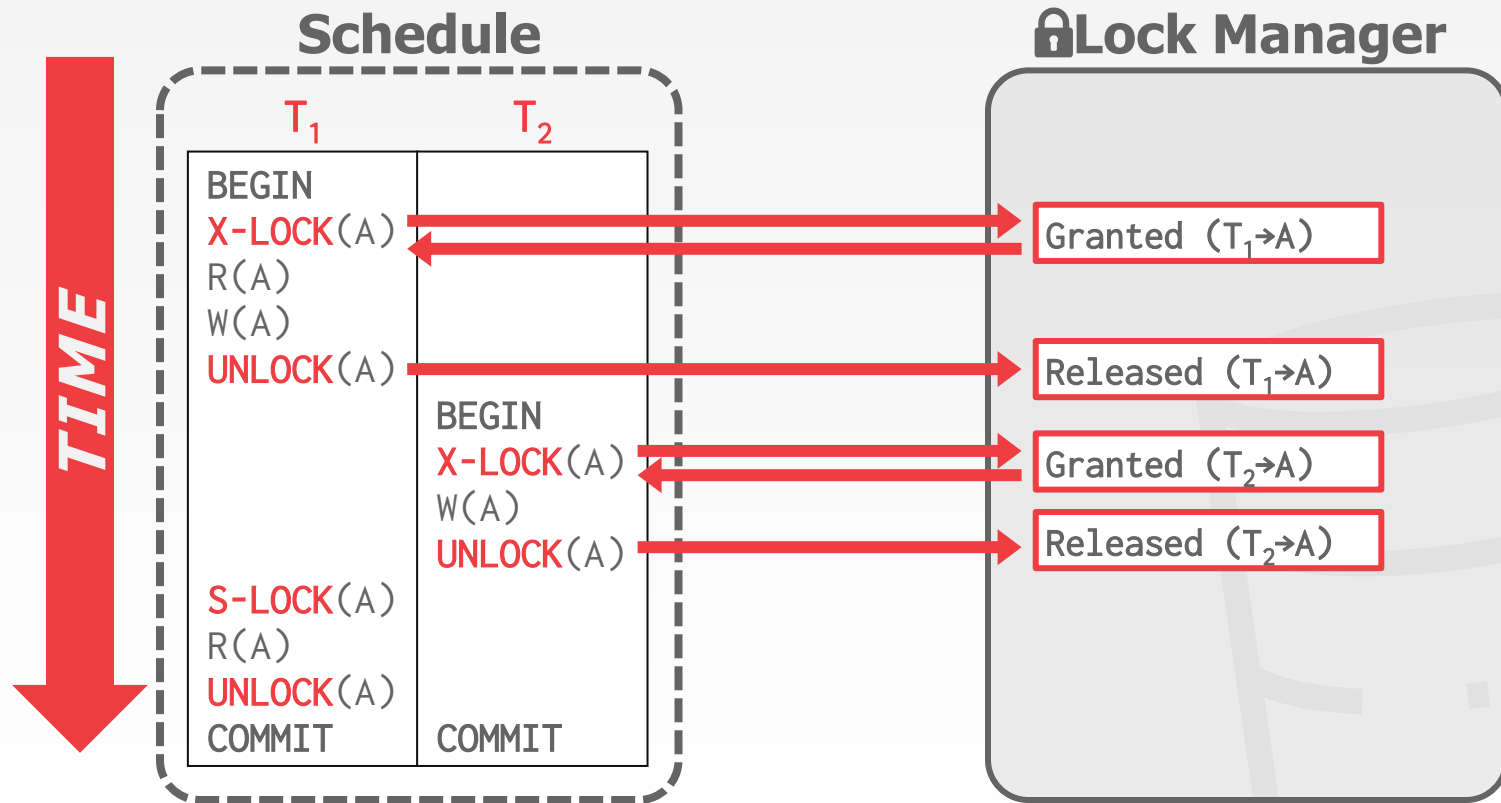
EXECUTING WITH LOCKS



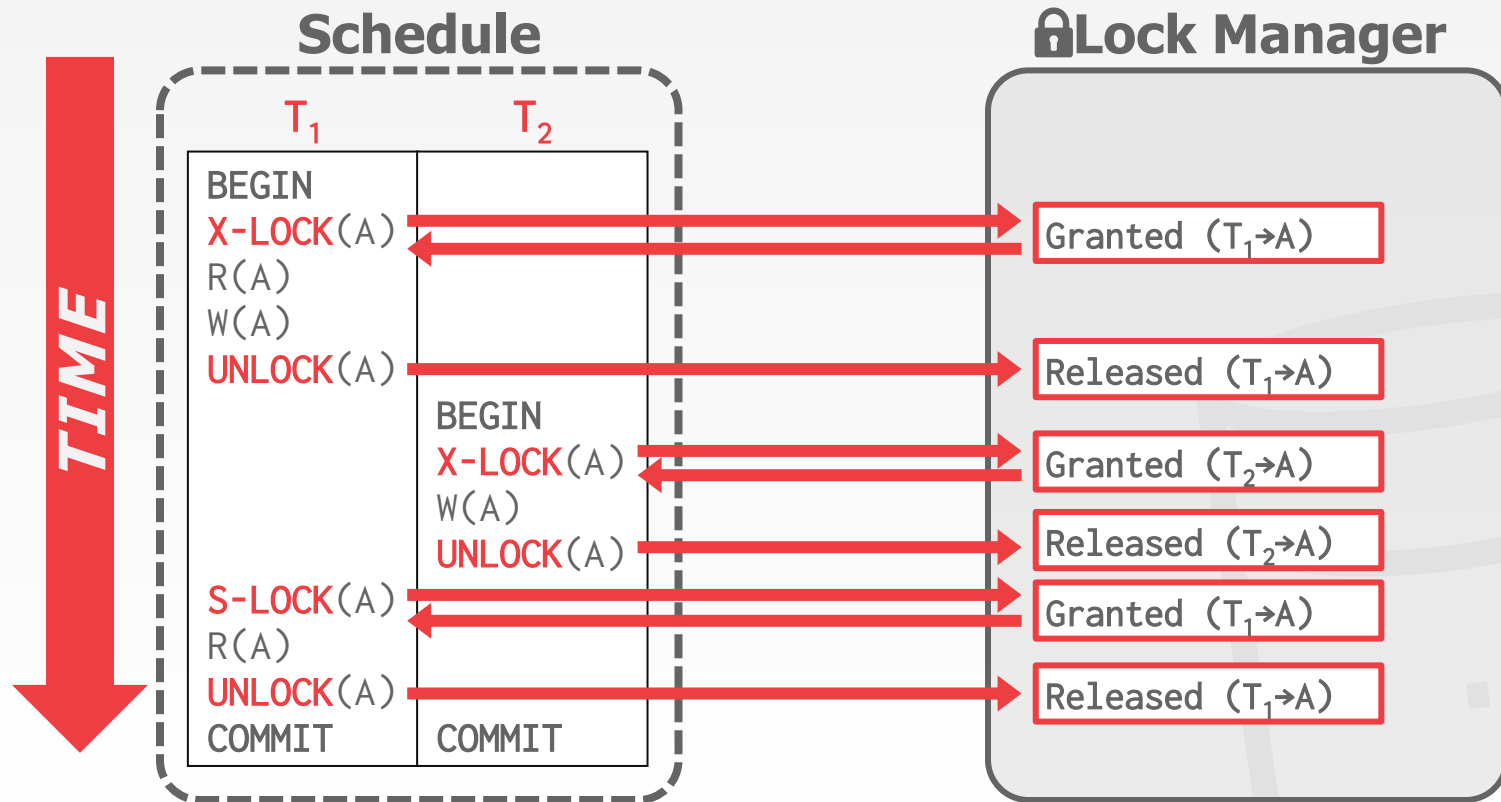
EXECUTING WITH LOCKS



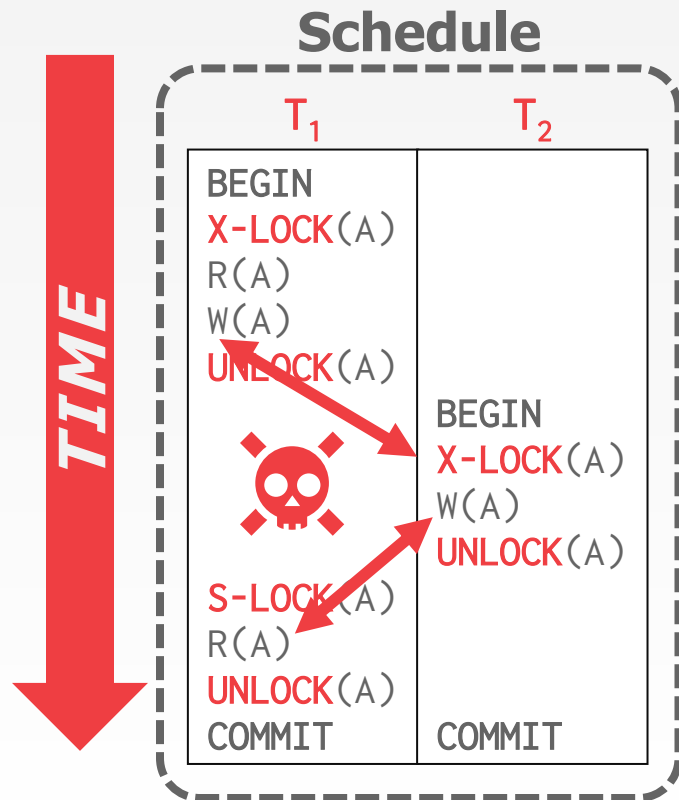
EXECUTING WITH LOCKS



EXECUTING WITH LOCKS



EXECUTING WITH LOCKS



Lock Manager

Granted ($T_1 \rightarrow A$)

Released ($T_1 \rightarrow A$)

Granted ($T_2 \rightarrow A$)

Released ($T_2 \rightarrow A$)

Granted ($T_1 \rightarrow A$)

Released ($T_1 \rightarrow A$)

CONCURRENCY CONTROL PROTOCOL

Two-phase locking (2PL) is a concurrency control protocol that determines whether a txn can access an object in the database on the fly.

The protocol does not need to know all the queries that a txn will execute ahead of time.



TWO-PHASE LOCKING

Phase #1: Growing

- Each txn requests the locks that it needs from the DBMS's lock manager.
- The lock manager grants/denies lock requests.

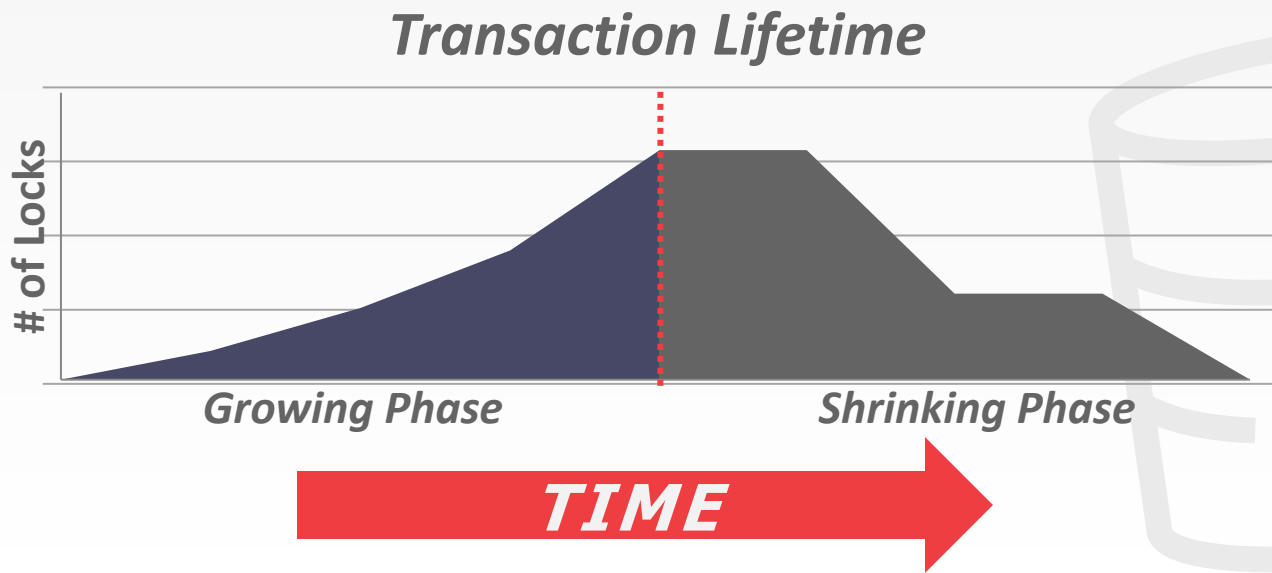
Phase #2: Shrinking

- The txn is allowed to only release locks that it previously acquired. It cannot acquire new locks.



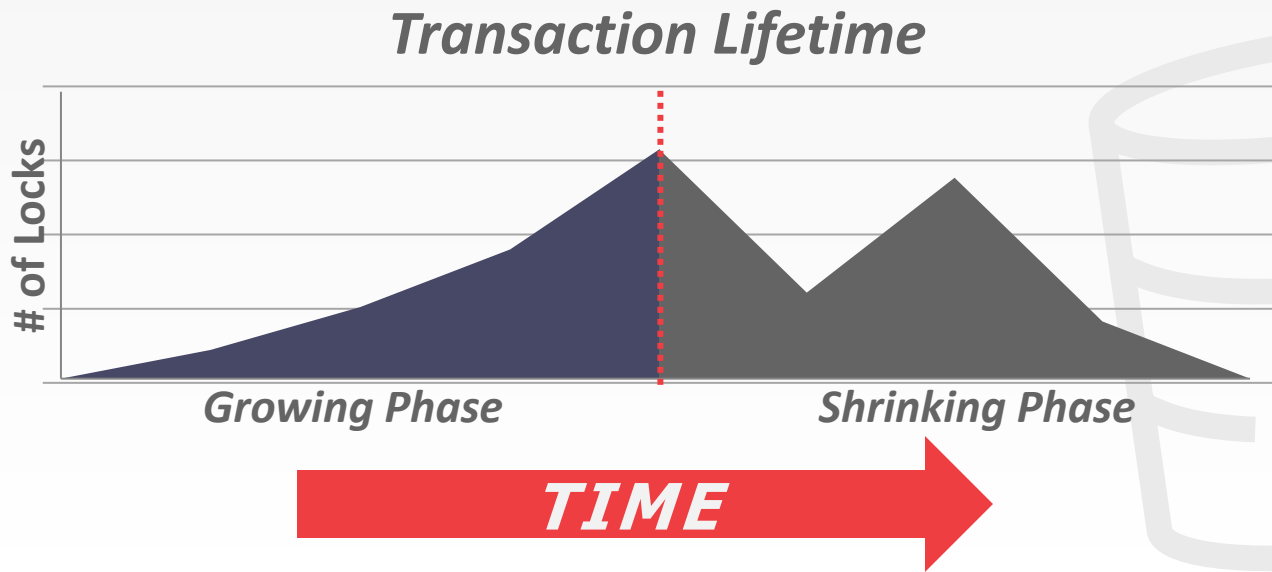
TWO-PHASE LOCKING

The txn is not allowed to acquire/upgrade locks after the growing phase finishes.



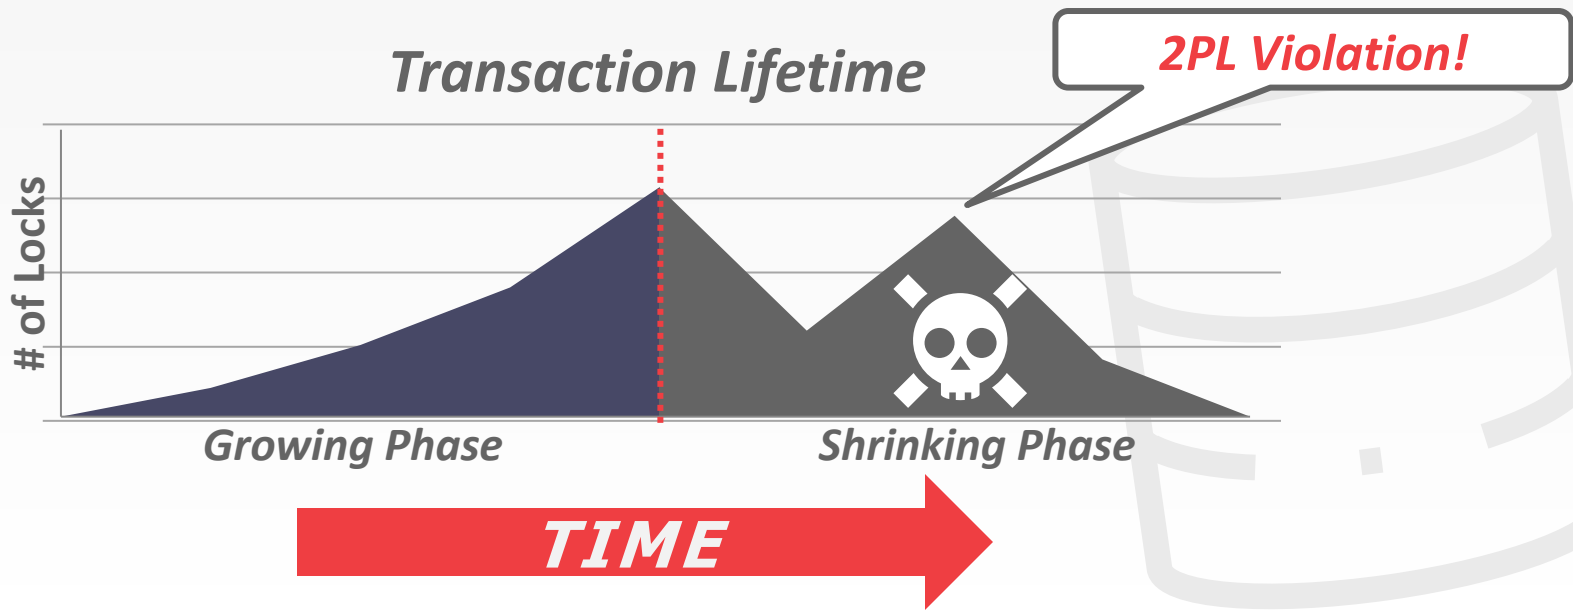
TWO-PHASE LOCKING

The txn is not allowed to acquire/upgrade locks after the growing phase finishes.

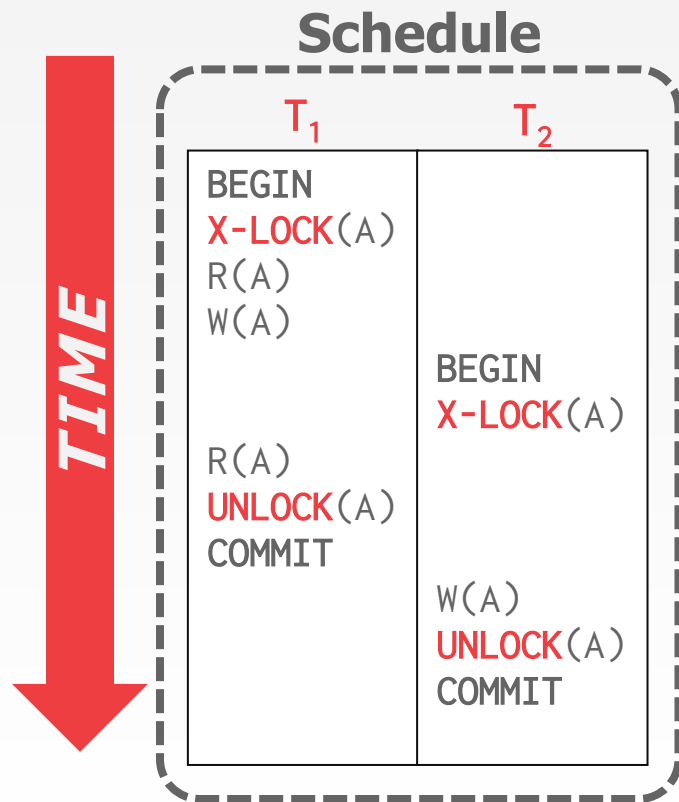


TWO-PHASE LOCKING

The txn is not allowed to acquire/upgrade locks after the growing phase finishes.



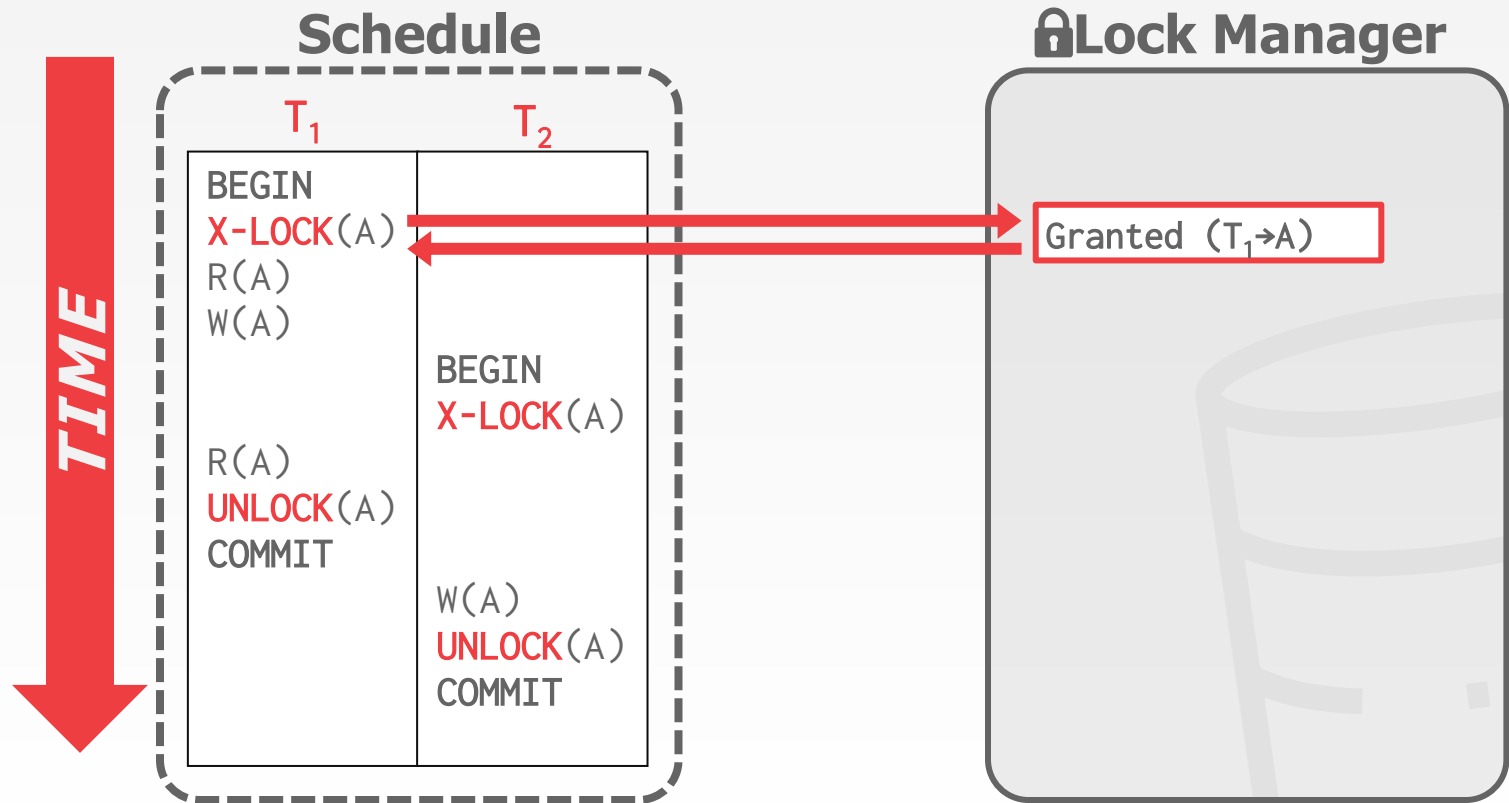
EXECUTING WITH 2PL



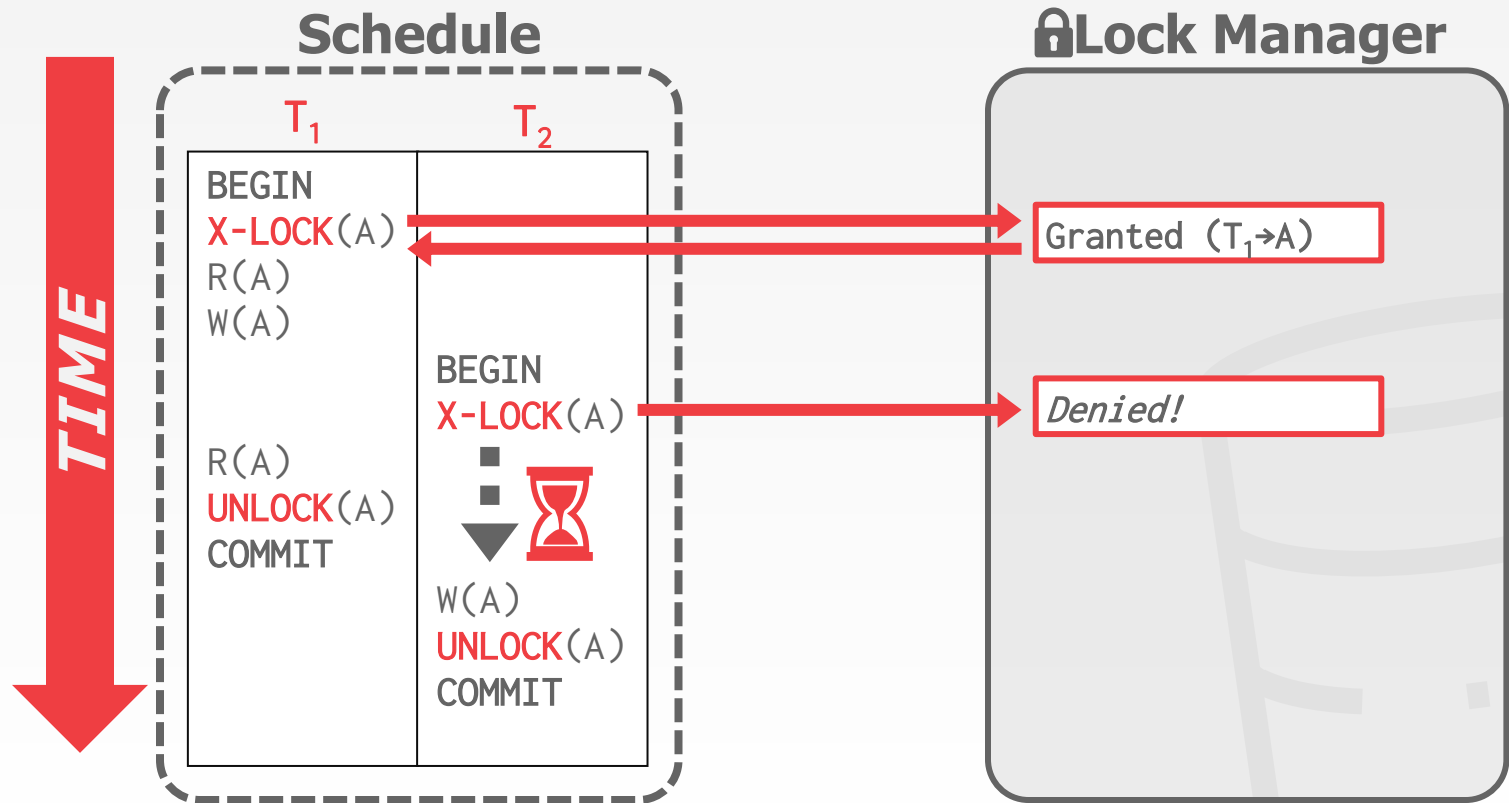
Lock Manager



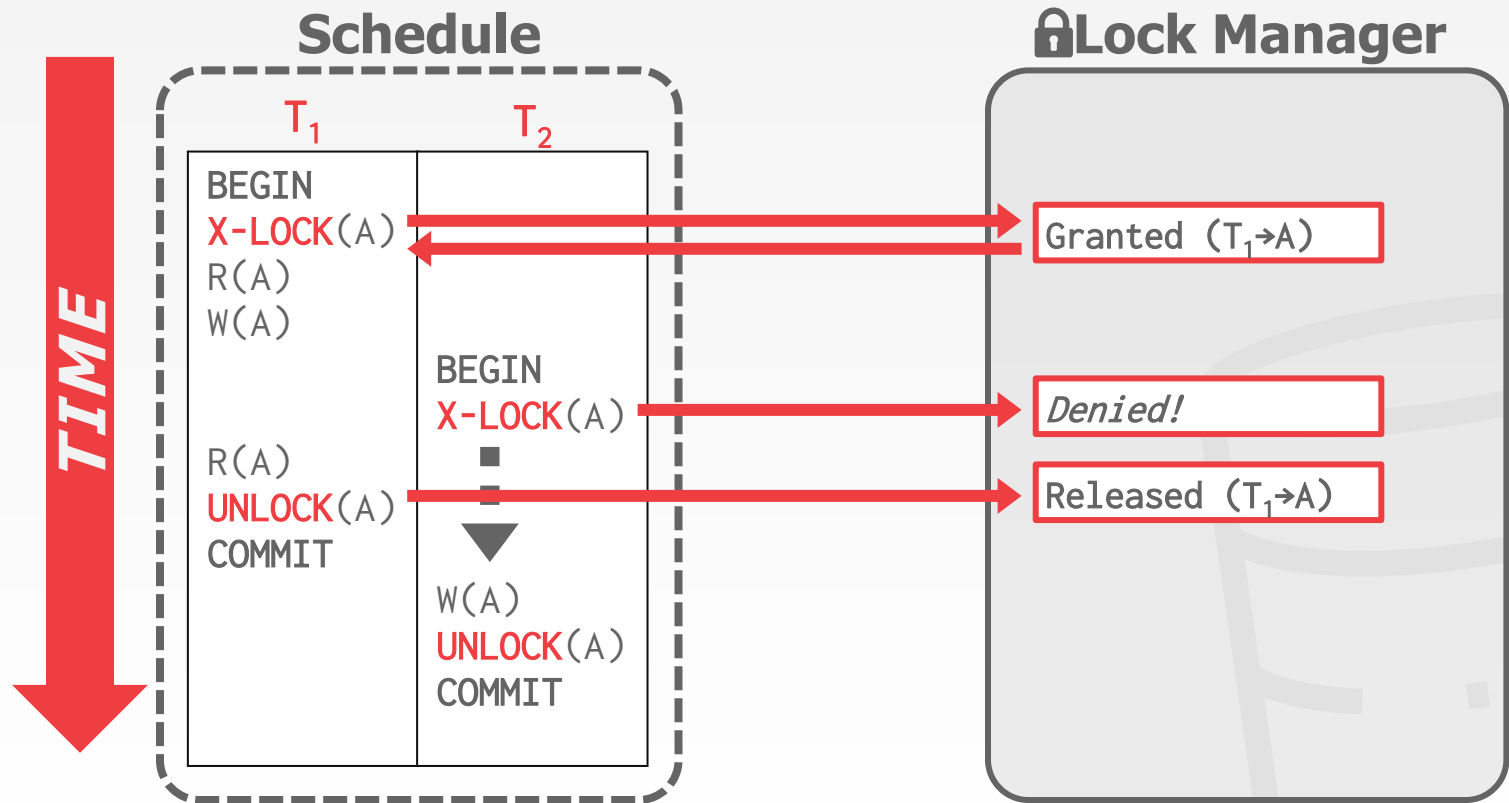
EXECUTING WITH 2PL



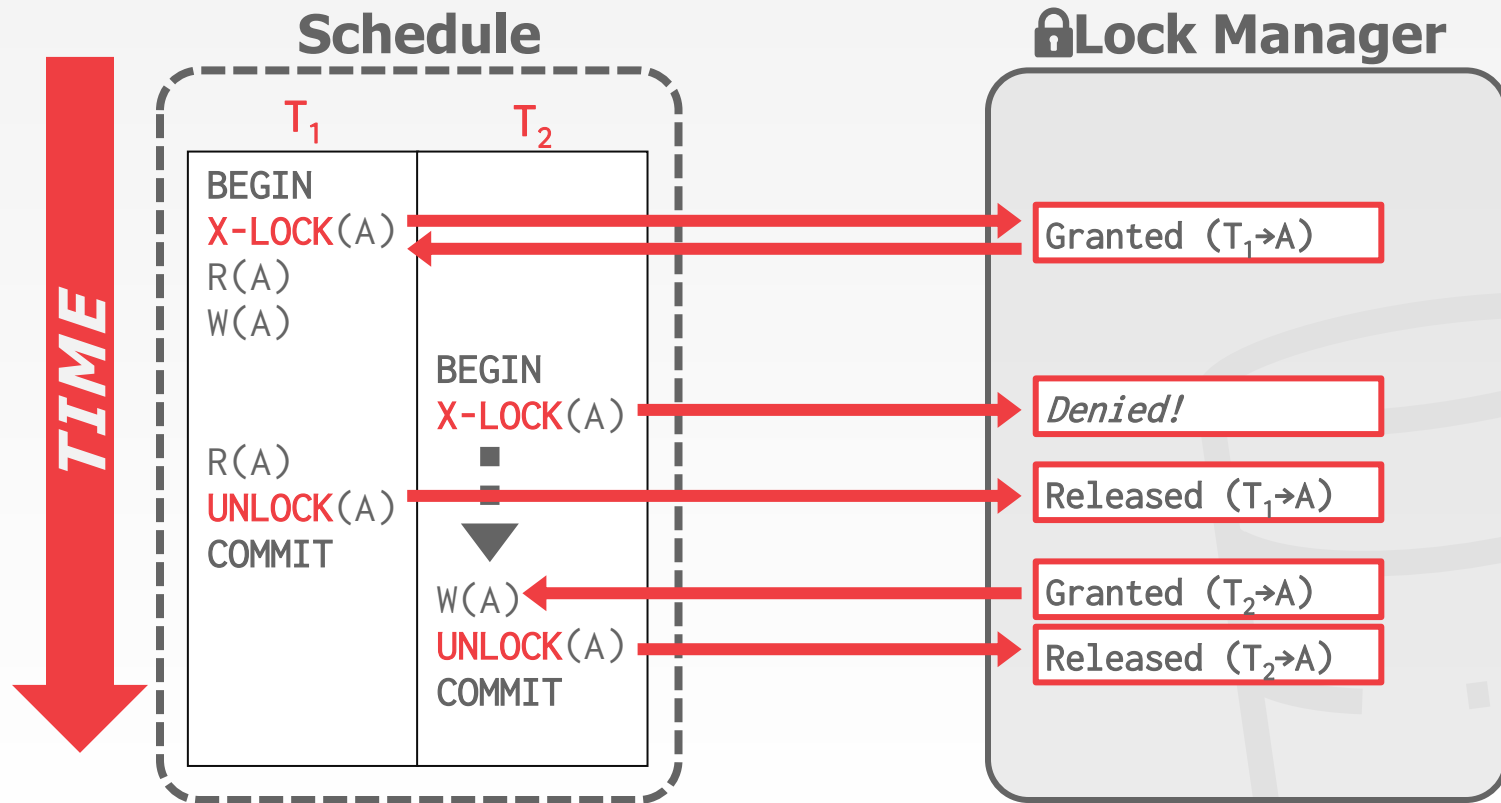
EXECUTING WITH 2PL



EXECUTING WITH 2PL



EXECUTING WITH 2PL



TWO-PHASE LOCKING

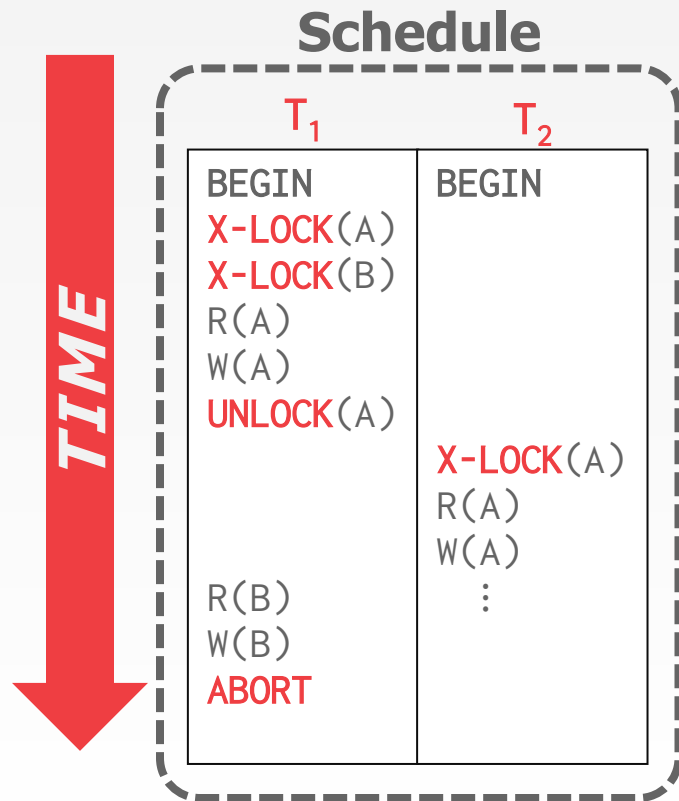
2PL on its own is sufficient to guarantee conflict serializability.

→ It generates schedules whose precedence graph is acyclic.

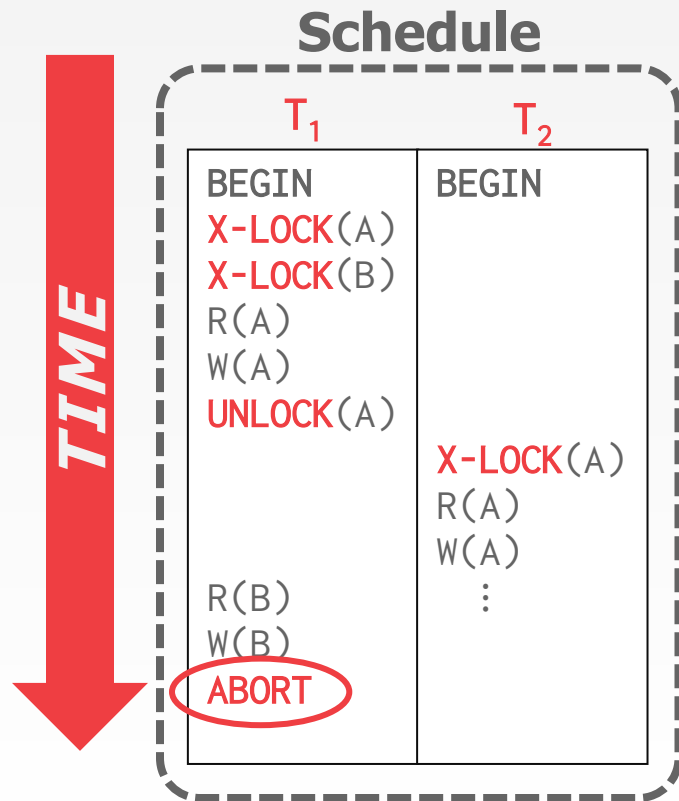
But it is subject to cascading aborts.



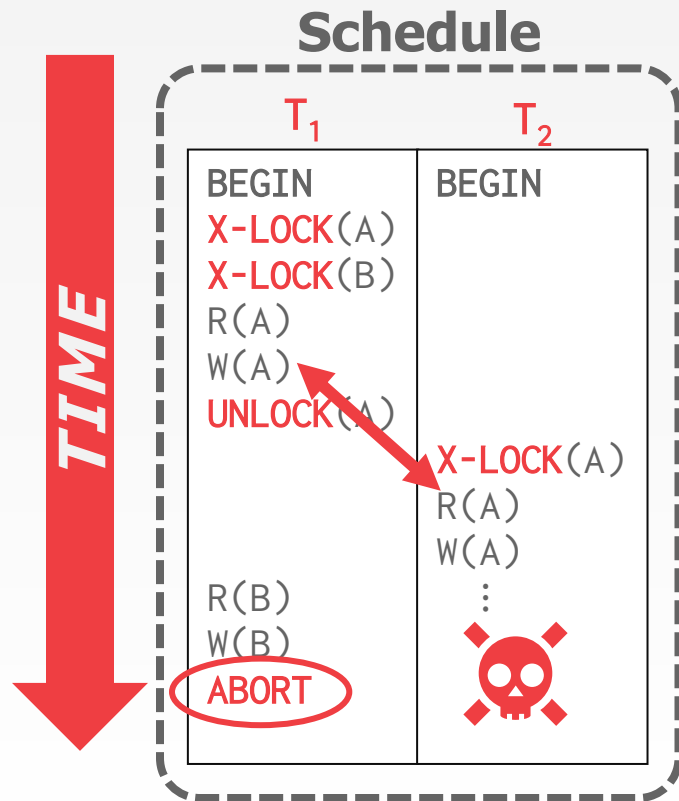
2PL – CASCADING ABORTS



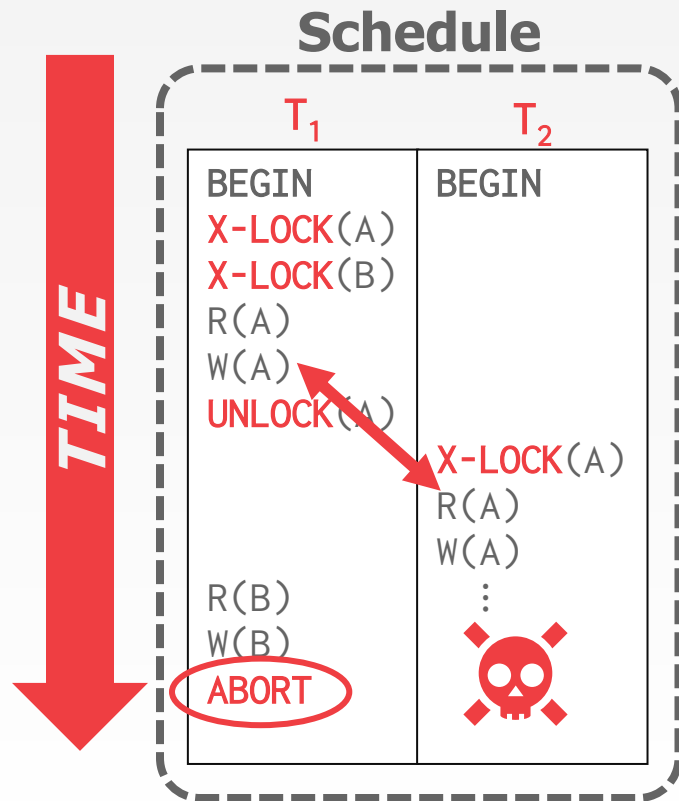
2PL – CASCADING ABORTS



2PL – CASCADING ABORTS



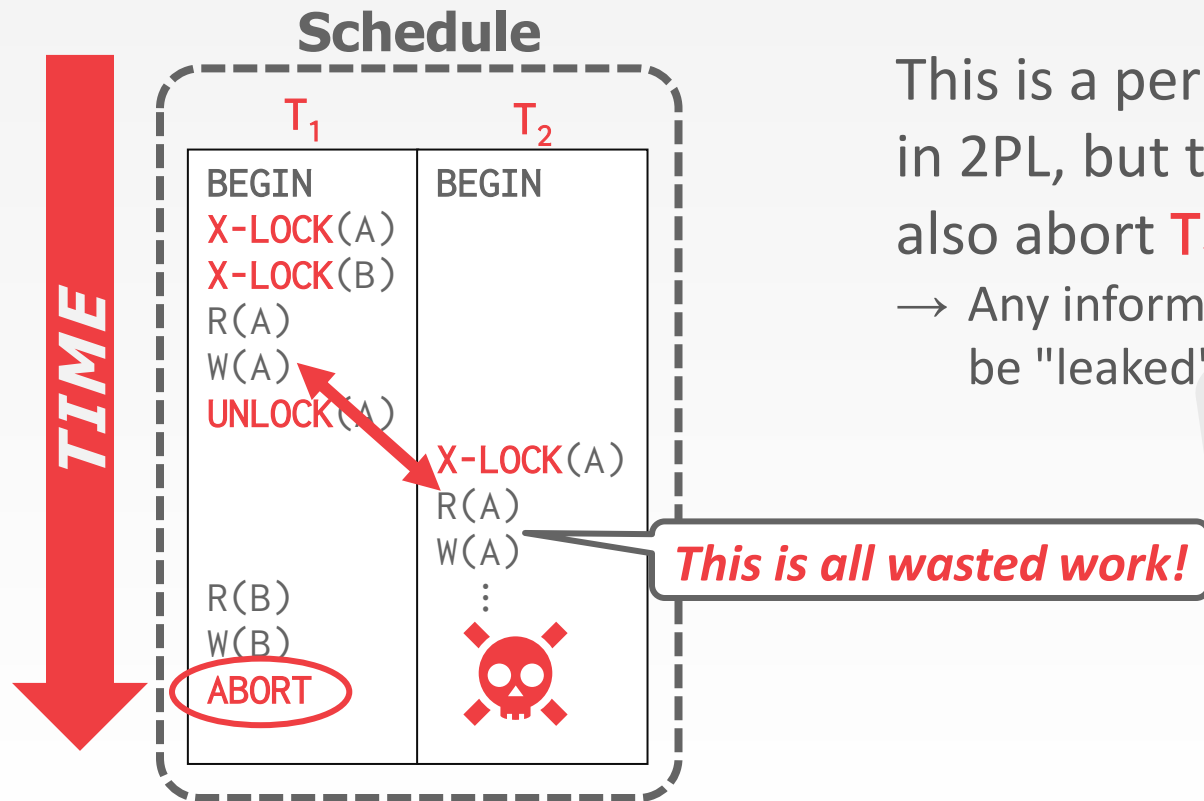
2PL – CASCADING ABORTS



This is a permissible schedule in 2PL, but the DBMS has to also abort T_2 when T_1 aborts.
 → Any information about T_1 cannot be "leaked" to the outside world.



2PL – CASCADING ABORTS



This is a permissible schedule in 2PL, but the DBMS has to also abort T_2 when T_1 aborts.
 → Any information about T_1 cannot be "leaked" to the outside world.

2PL OBSERVATIONS

There are potential schedules that are serializable but would not be allowed by 2PL.

→ Locking limits concurrency.

May still have "dirty reads".

→ Solution: **Strong Strict 2PL (aka Rigorous 2PL)**

May lead to deadlocks.

→ Solution: **Detection or Prevention**



2PL OBSERVATIONS

There are potential schedules that are serializable but would not be allowed by 2PL.

→ Locking limits concurrency.

May still have "dirty reads".

→ Solution: **Strong Strict 2PL (aka Rigorous 2PL)**

May lead to deadlocks.

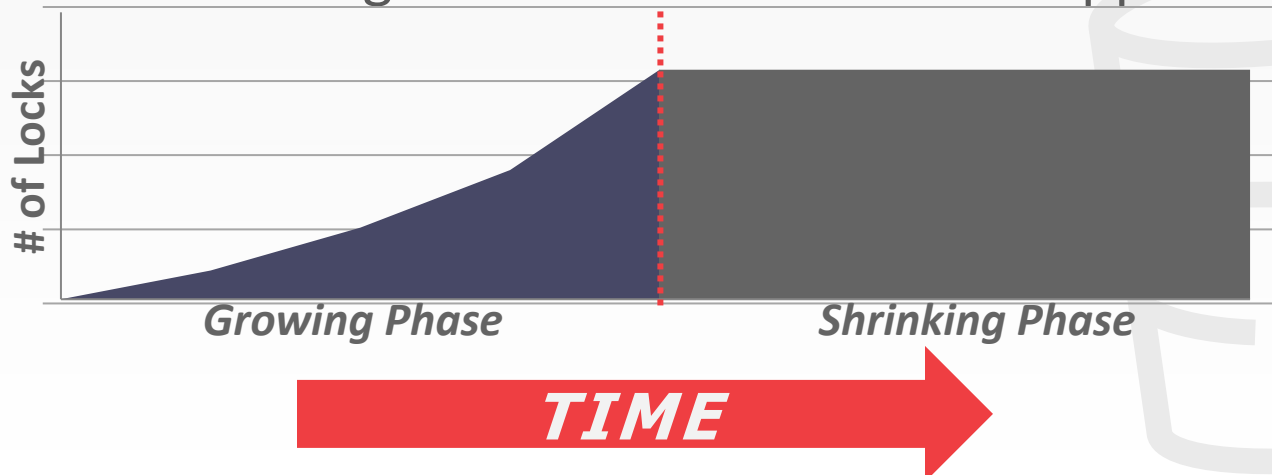
→ Solution: **Detection or Prevention**



STRONG STRICT TWO-PHASE LOCKING

The txn is only allowed to release locks after it has ended, i.e., committed or aborted.

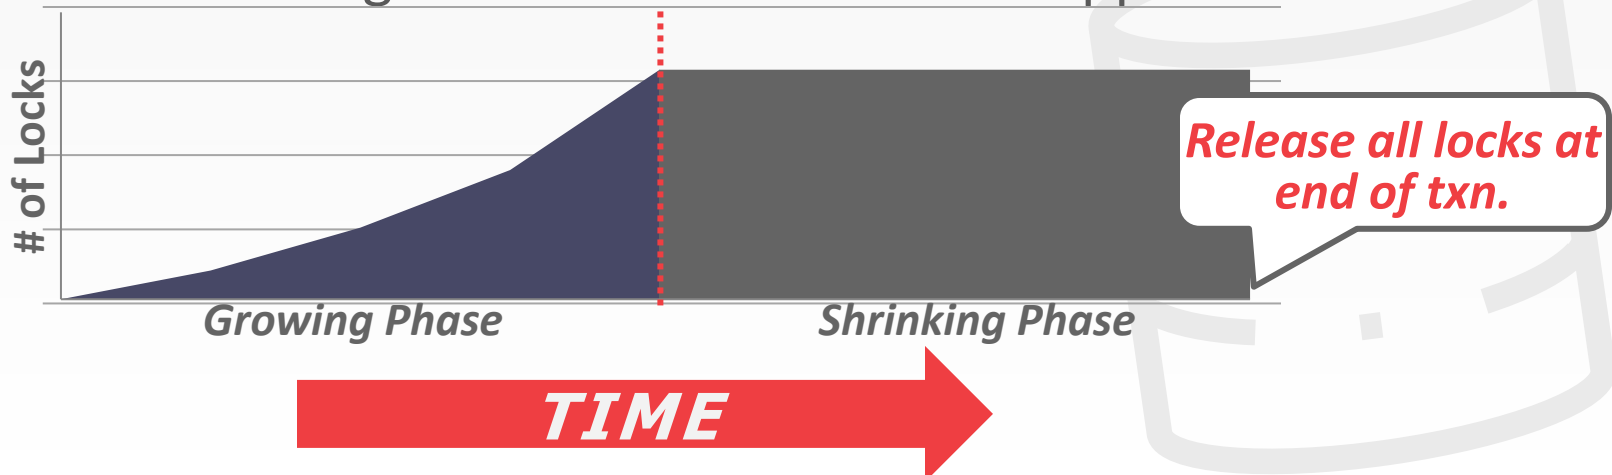
Allows only conflict serializable schedules, but it is often stronger than needed for some apps.



STRONG STRICT TWO-PHASE LOCKING

The txn is only allowed to release locks after it has ended, i.e., committed or aborted.

Allows only conflict serializable schedules, but it is often stronger than needed for some apps.



STRONG STRICT TWO-PHASE LOCKING

A schedule is **strict** if a value written by a txn is not read or overwritten by other txns until that txn finishes.

Advantages:

- Does not incur cascading aborts.
- Aborted txns can be undone by just restoring original values of modified tuples.



EXAMPLES

T_1 – Move \$100 from Lin's account (**A**) to his friend's account (**B**).

T_2 – Compute the total amount in all accounts and return it to the application.

T_1

```
BEGIN  
A=A-100  
B=B+100  
COMMIT
```

T_2

```
BEGIN  
ECHO A+B  
COMMIT
```



EXAMPLES

T_1 – Move \$100 from Lin's account (**A**) to his friend's account (**B**).

T_2 – Compute the total amount in all accounts and return it to the application.

T_1

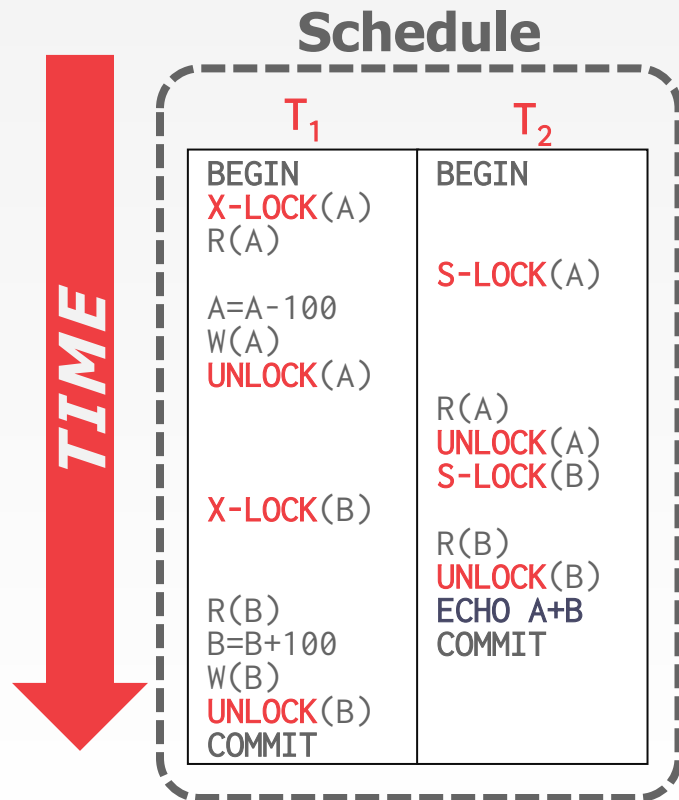
```
BEGIN
A=A-100
B=B+100
COMMIT
```

T_2

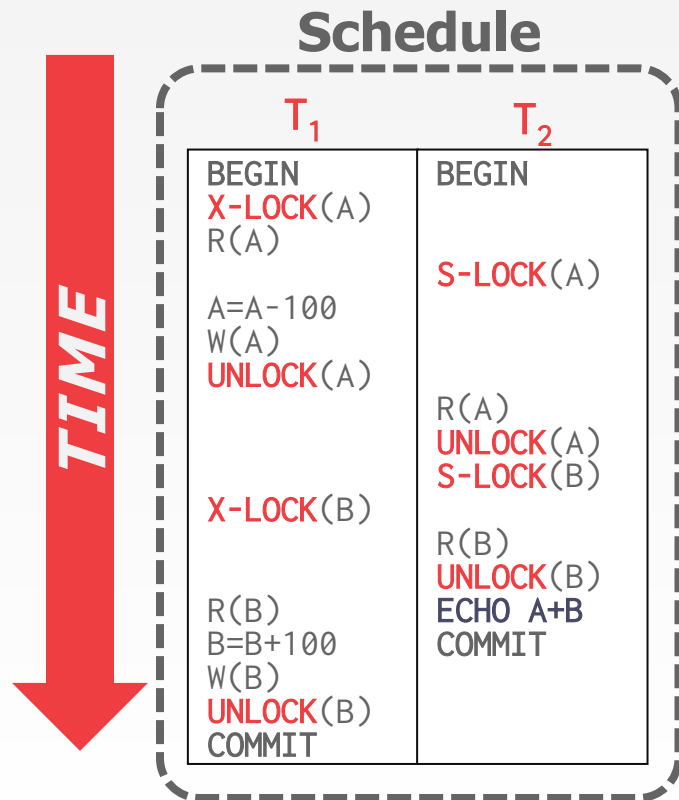
```
BEGIN
ECHO A+B
COMMIT
```



NON-2PL EXAMPLE



NON-2PL EXAMPLE

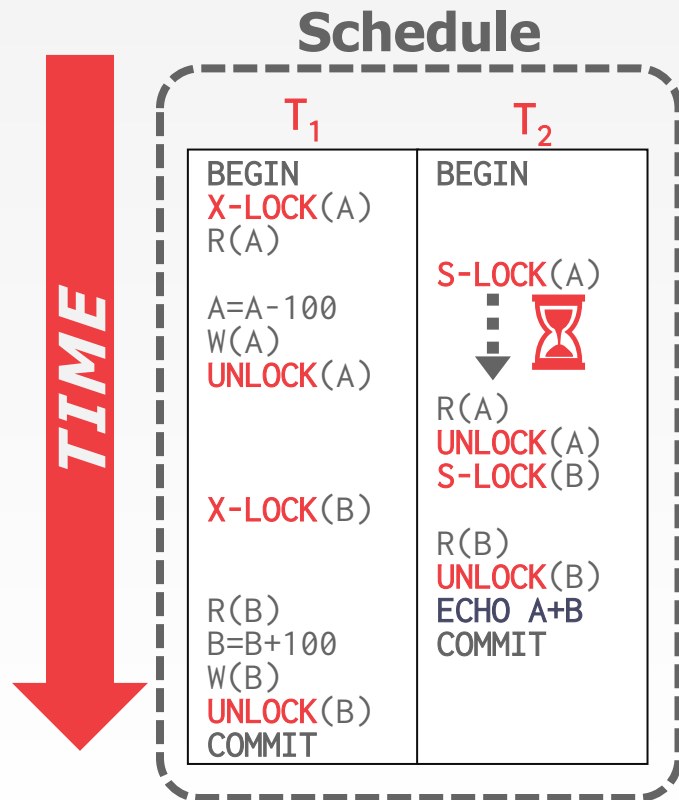


Initial Database State

A=1000, **B**=1000



NON-2PL EXAMPLE

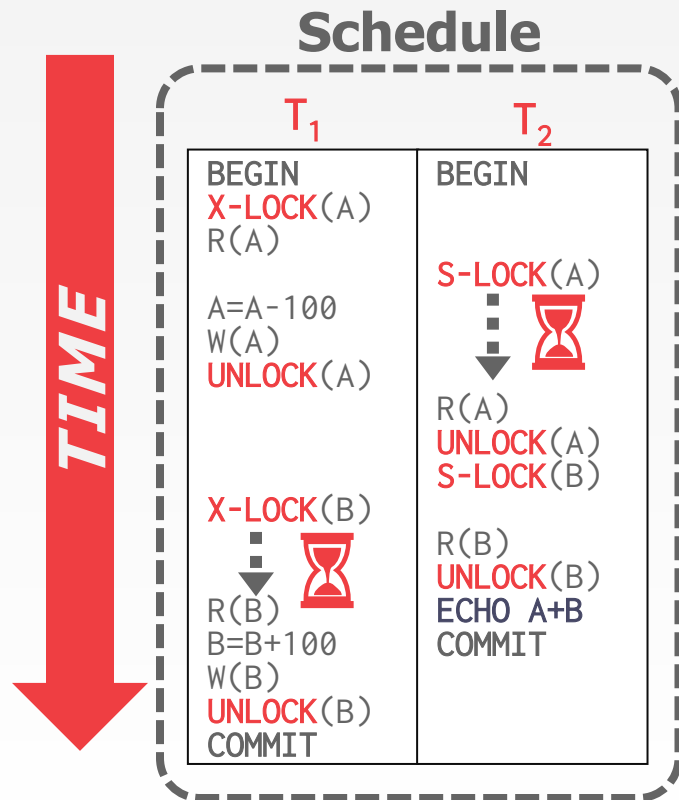


Initial Database State

A=1000, **B**=1000



NON-2PL EXAMPLE

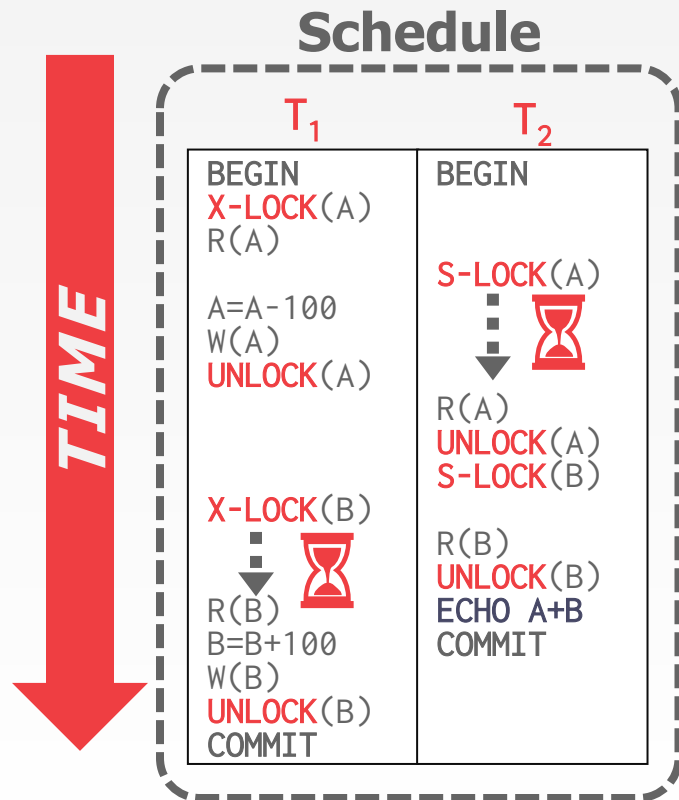


Initial Database State

A=1000, **B**=1000



NON-2PL EXAMPLE



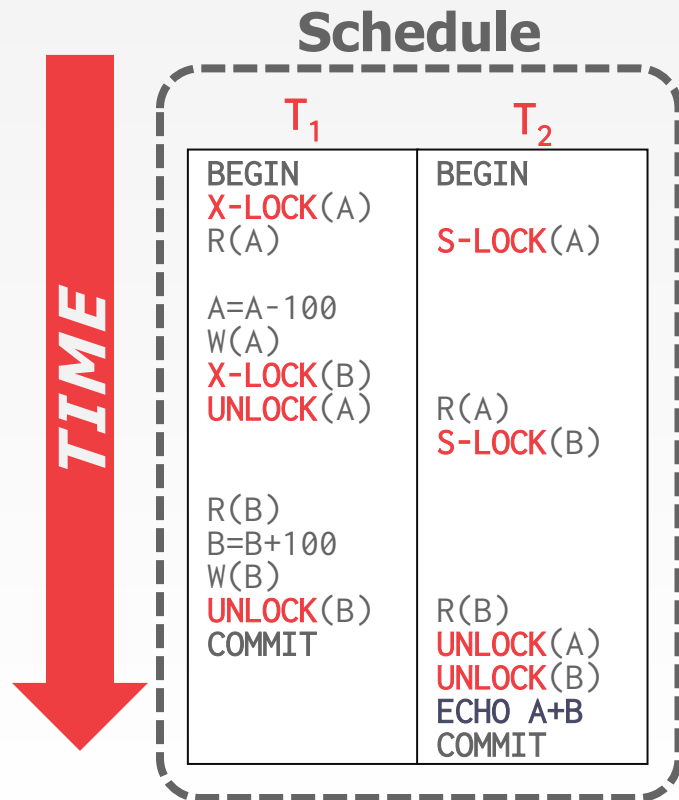
Initial Database State

A=1000, **B**=1000

T_2 Output

A+B=1900

2PL EXAMPLE

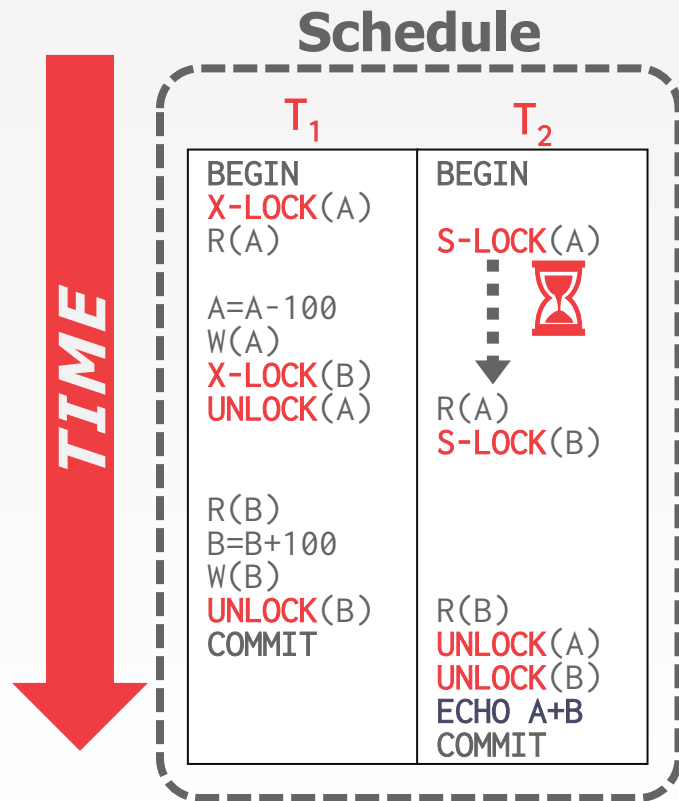


Initial Database State

A=1000, **B**=1000



2PL EXAMPLE

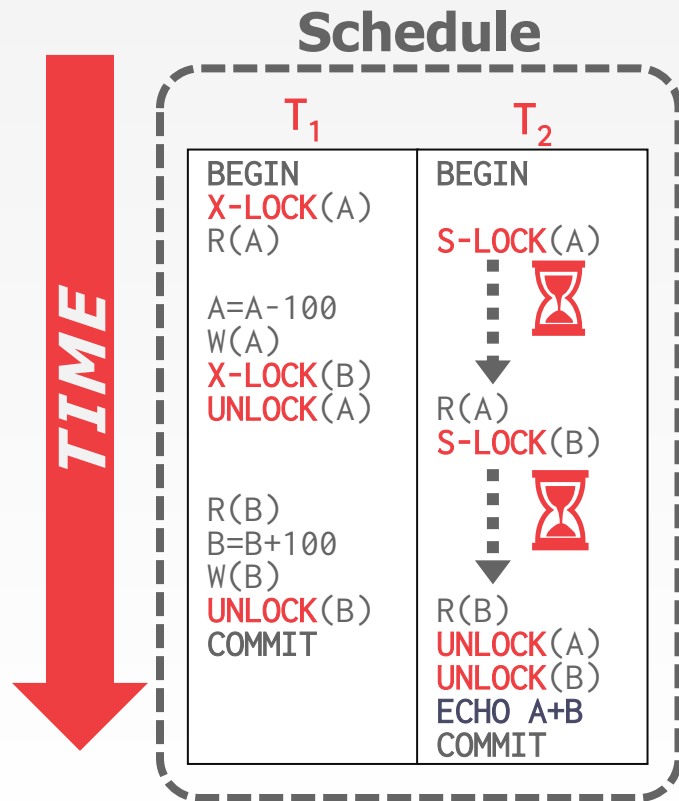


Initial Database State

A=1000, **B**=1000



2PL EXAMPLE

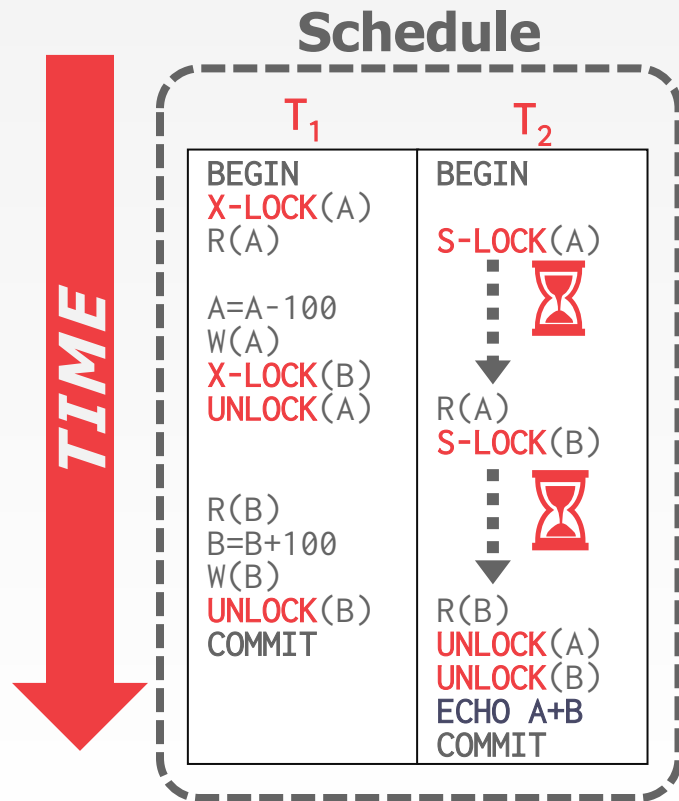


Initial Database State

A=1000, **B**=1000



2PL EXAMPLE



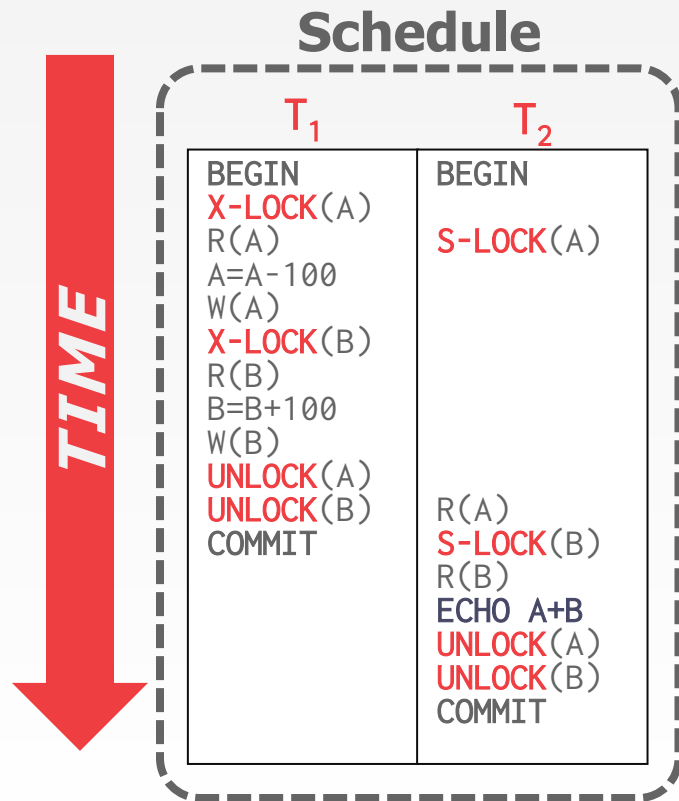
Initial Database State

A=1000, **B**=1000

T_2 Output

A+**B**=2000

STRONG STRICT 2PL EXAMPLE

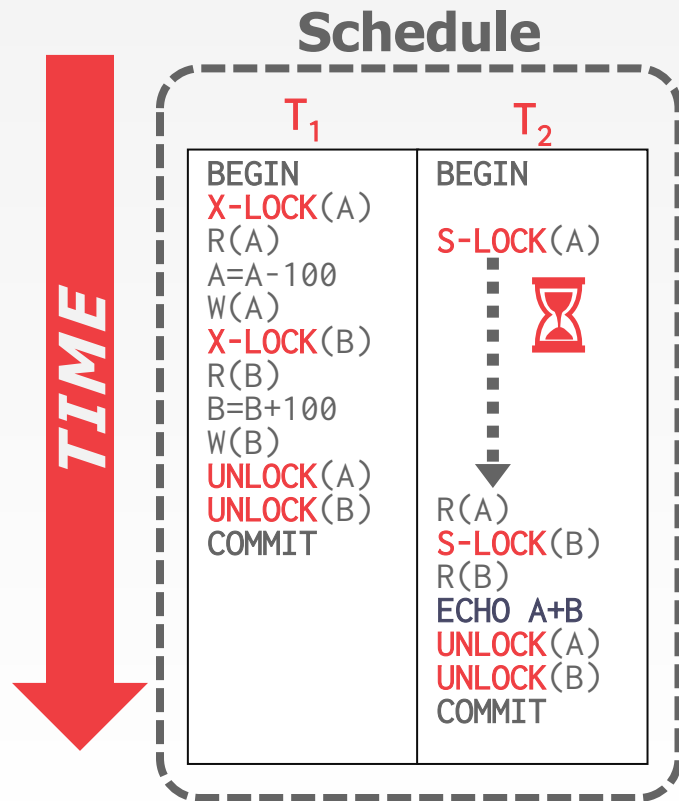


Initial Database State

A=1000, **B**=1000



STRONG STRICT 2PL EXAMPLE

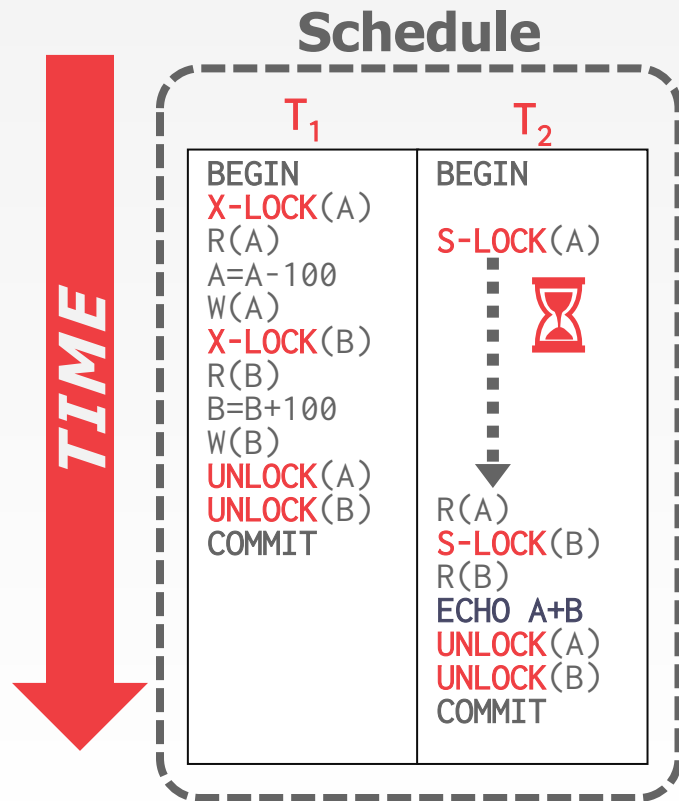


Initial Database State

A=1000, **B**=1000



STRONG STRICT 2PL EXAMPLE



Initial Database State

A=1000, **B**=1000

T_2 Output

A+B=2000

UNIVERSE OF SCHEDULES

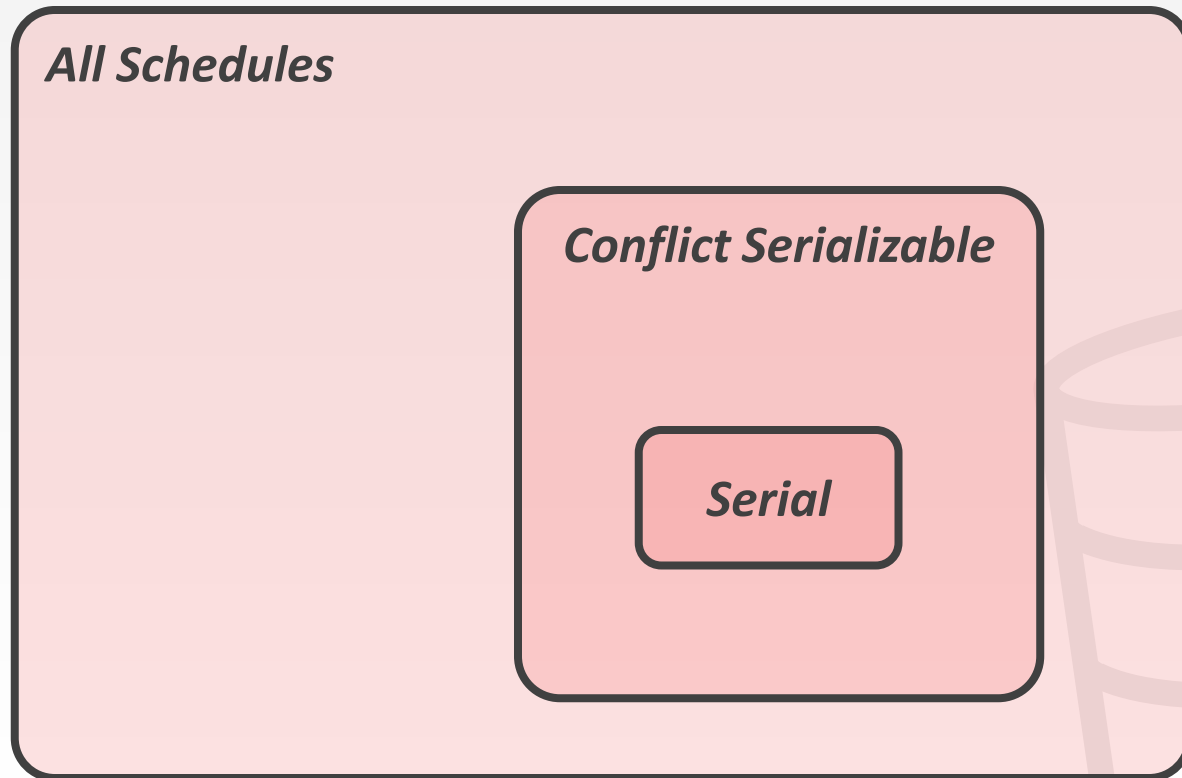
All Schedules

UNIVERSE OF SCHEDULES

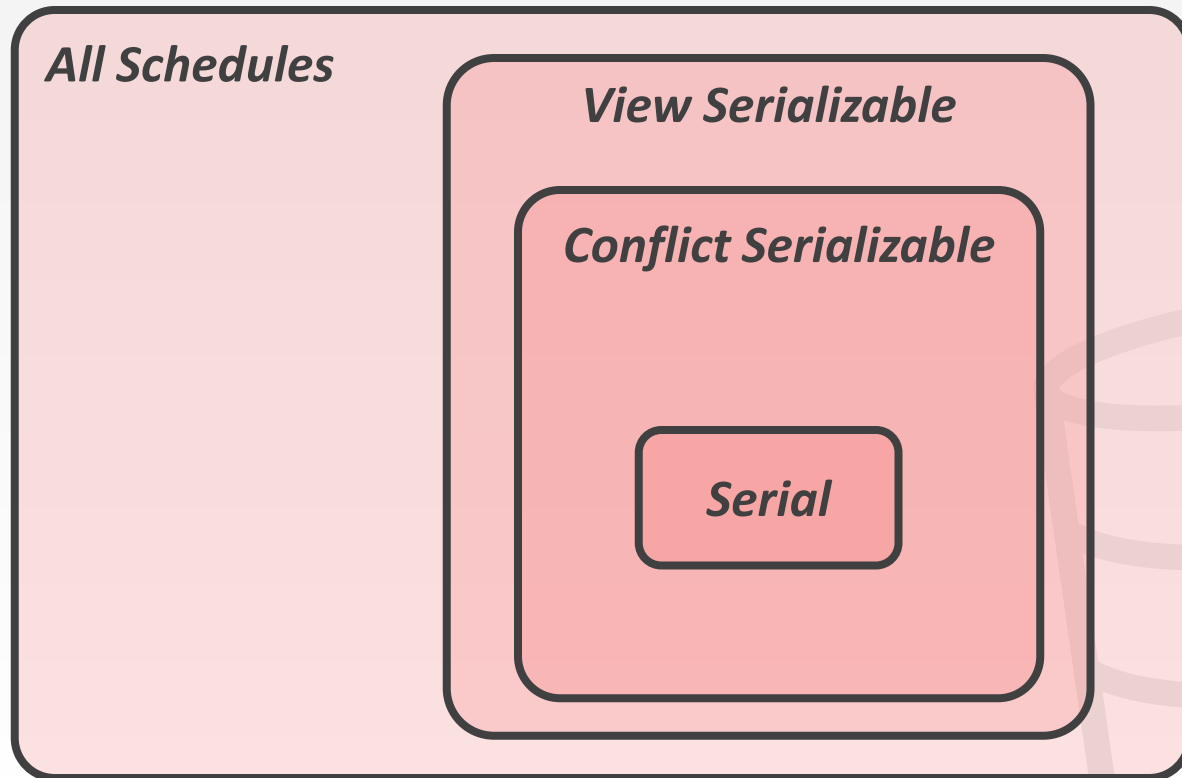
All Schedules

Serial

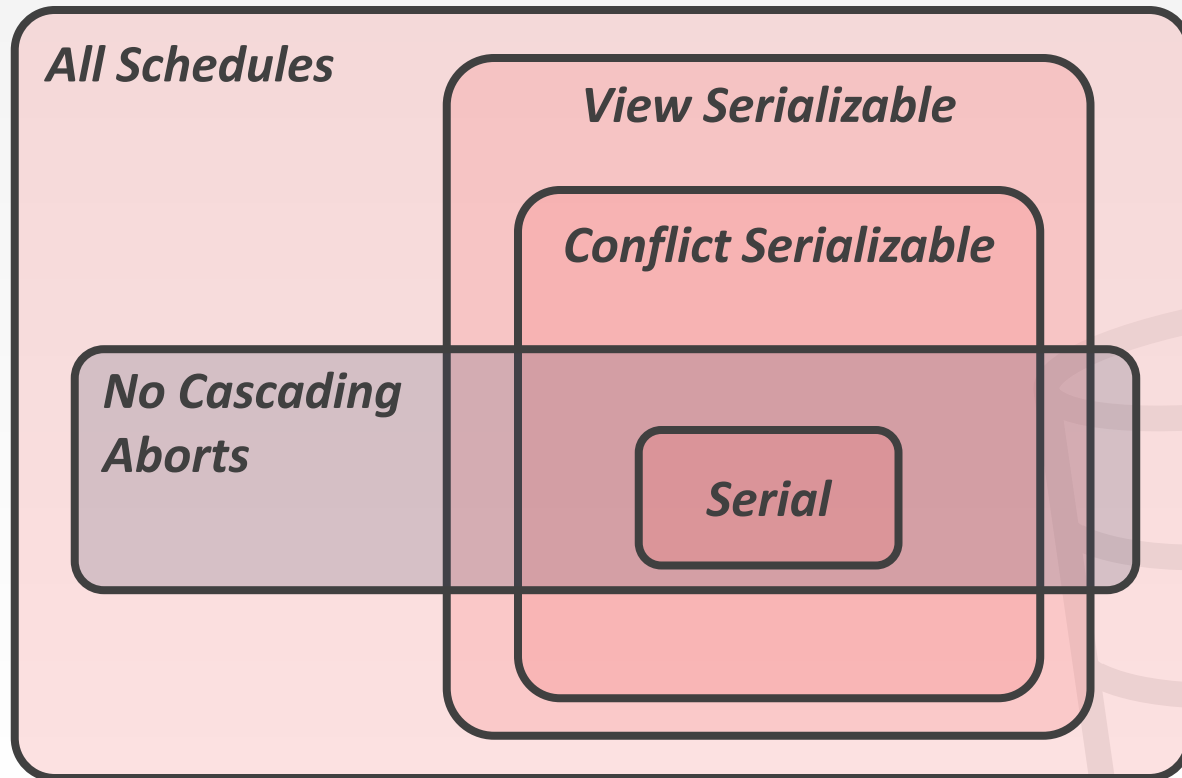
UNIVERSE OF SCHEDULES



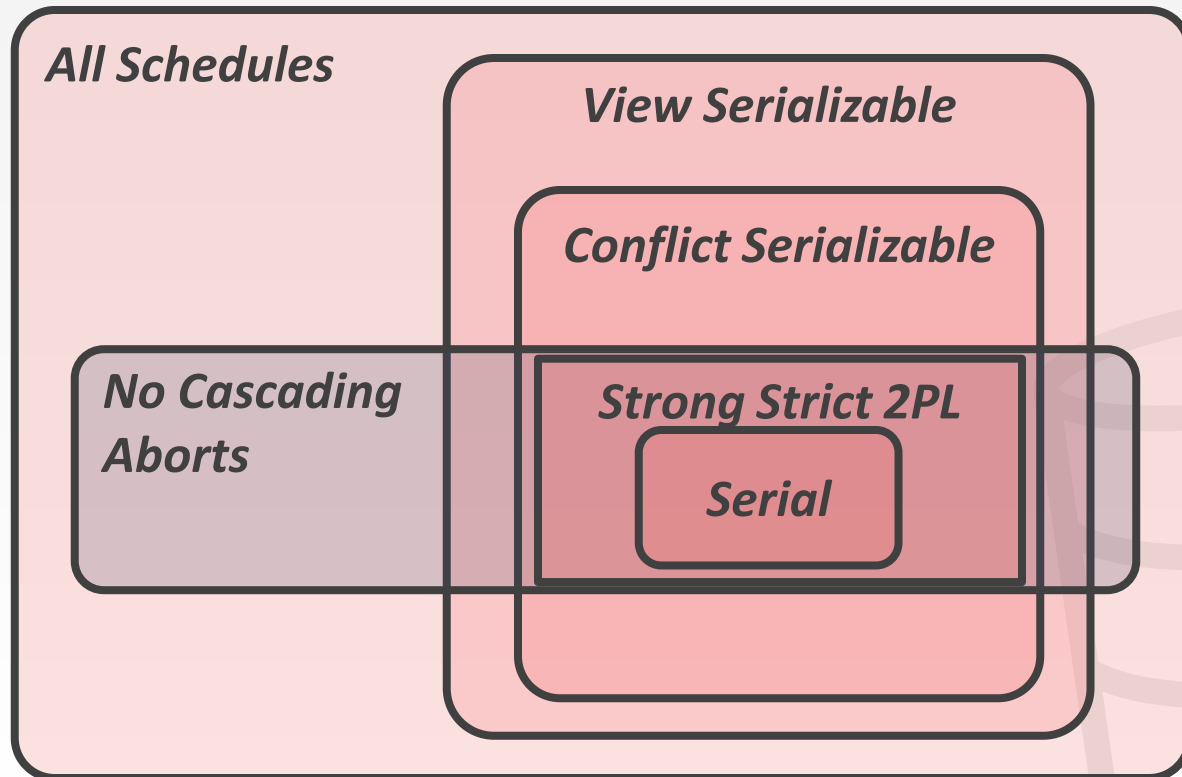
UNIVERSE OF SCHEDULES



UNIVERSE OF SCHEDULES



UNIVERSE OF SCHEDULES



2PL OBSERVATIONS

There are potential schedules that are serializable but would not be allowed by 2PL.

→ Locking limits concurrency.

May still have "dirty reads".

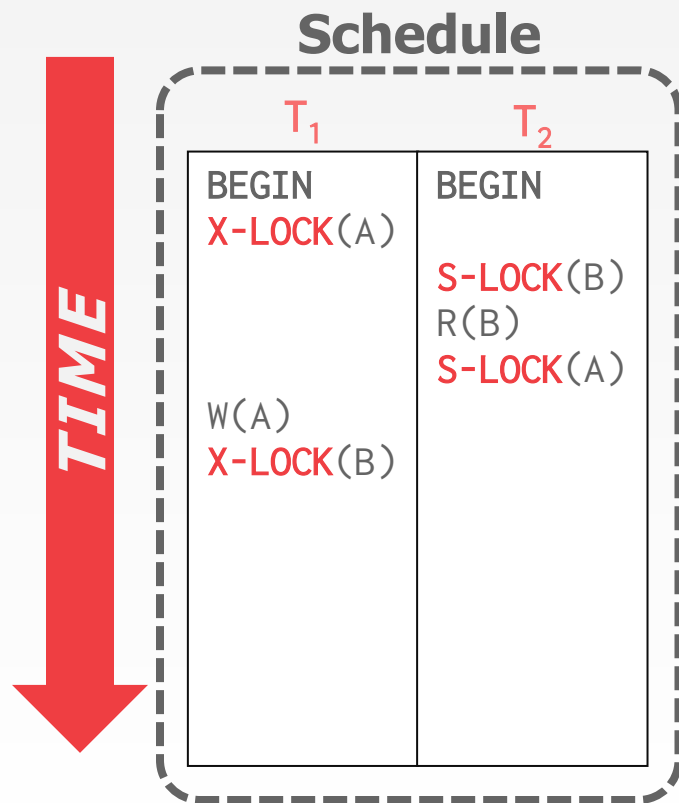
→ Solution: **Strong Strict 2PL (Rigorous)**

May lead to deadlocks.

→ Solution: **Detection or Prevention**

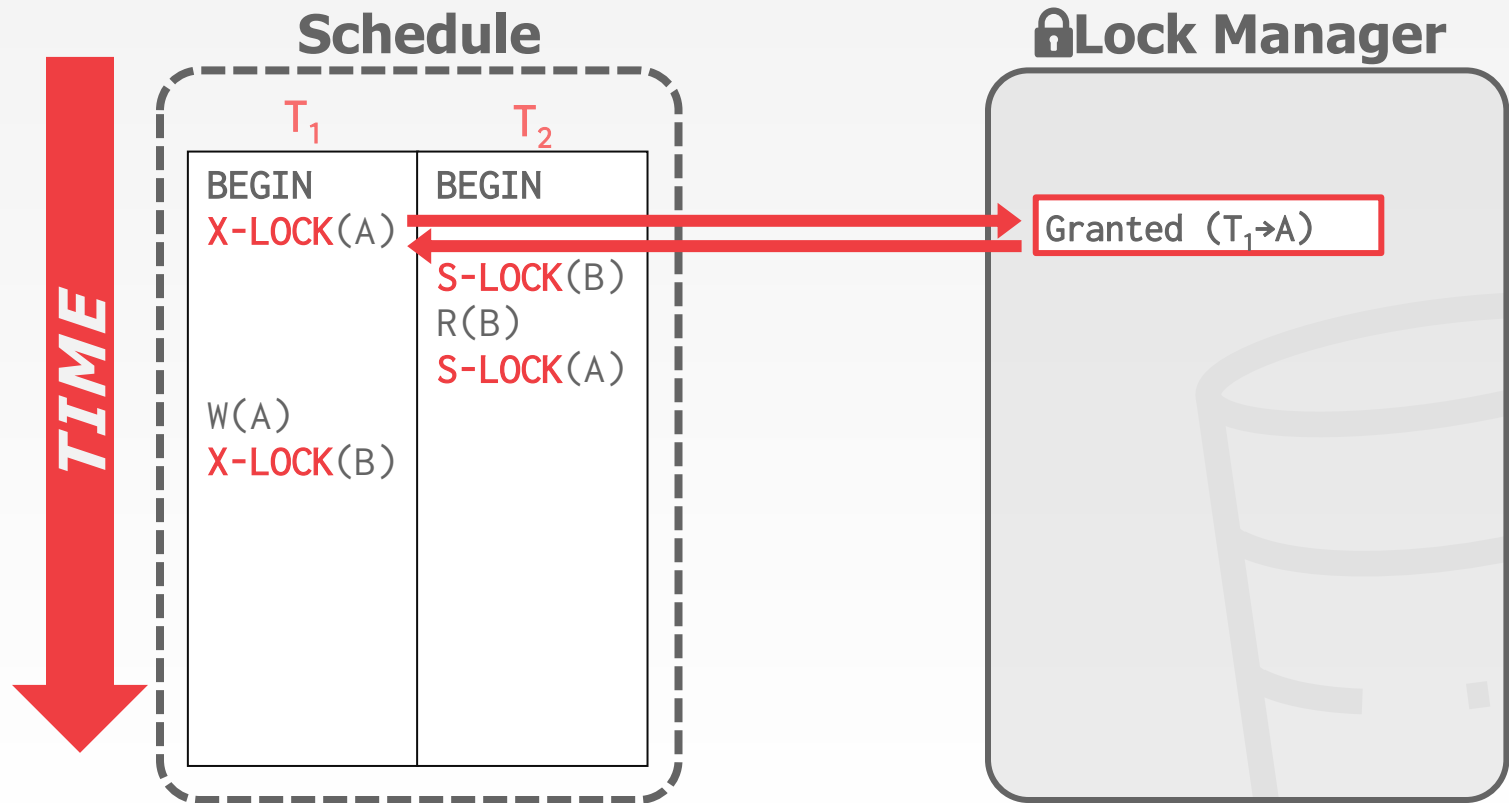


2PL DEADLOCKS

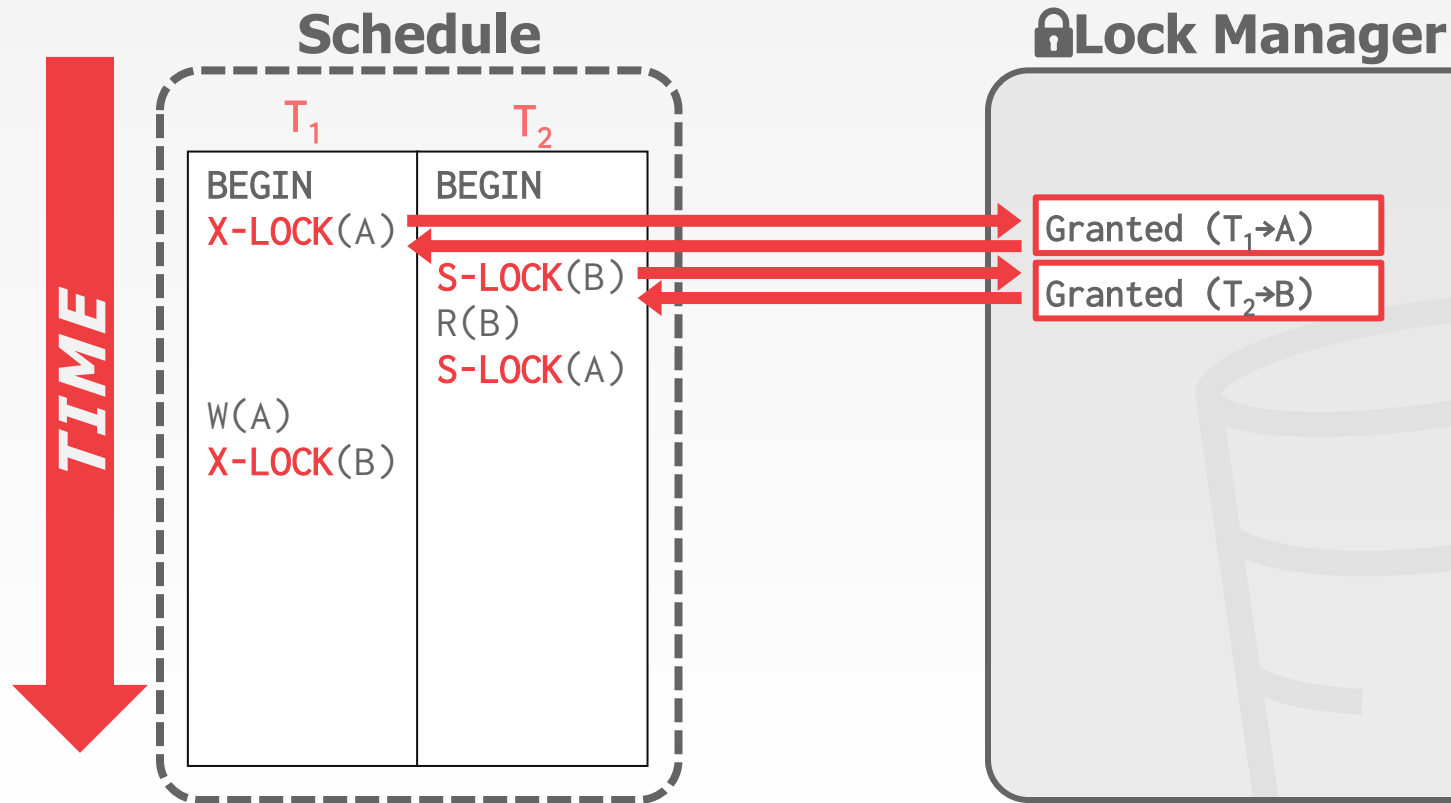


Lock Manager

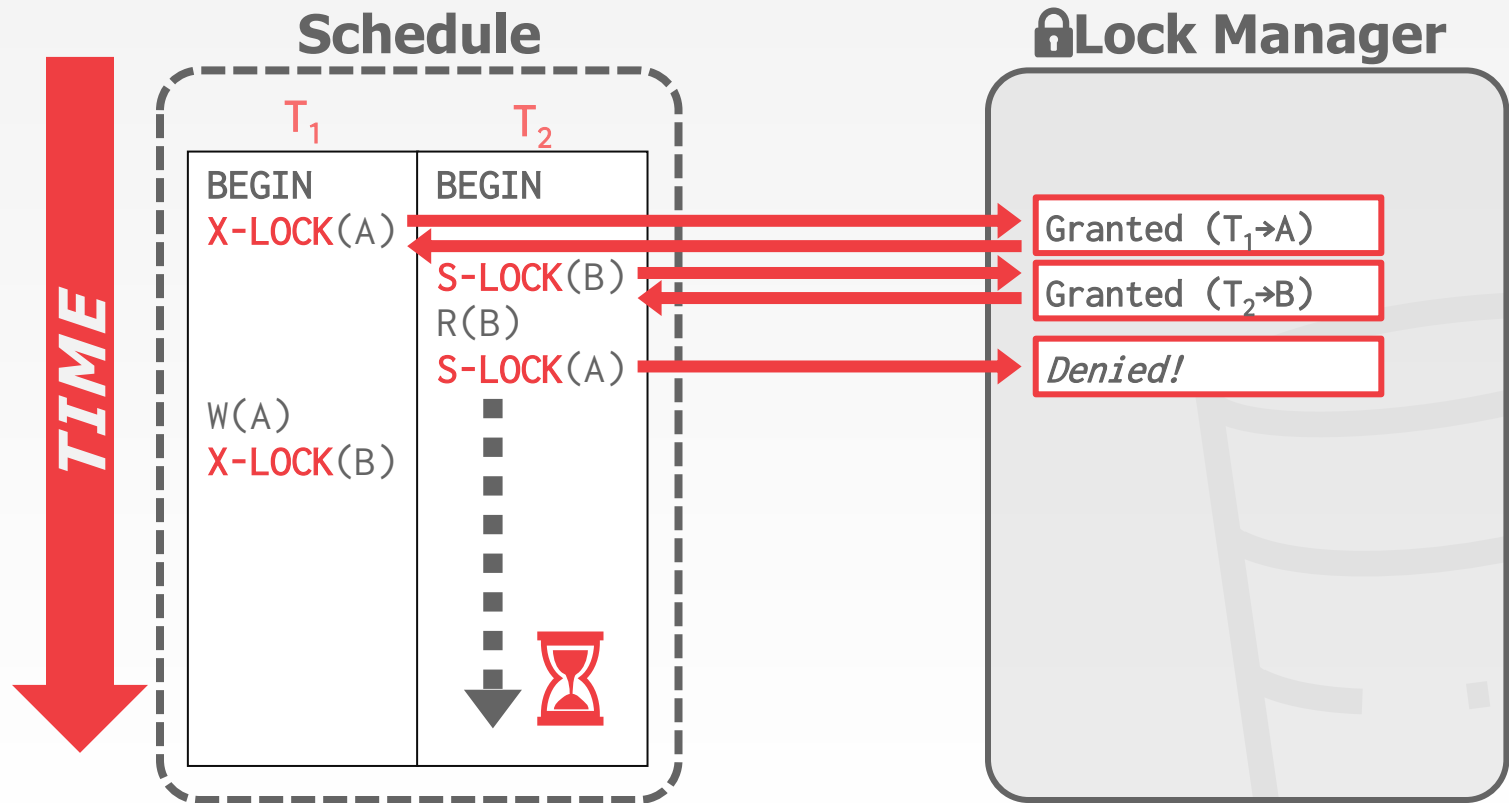
2PL DEADLOCKS



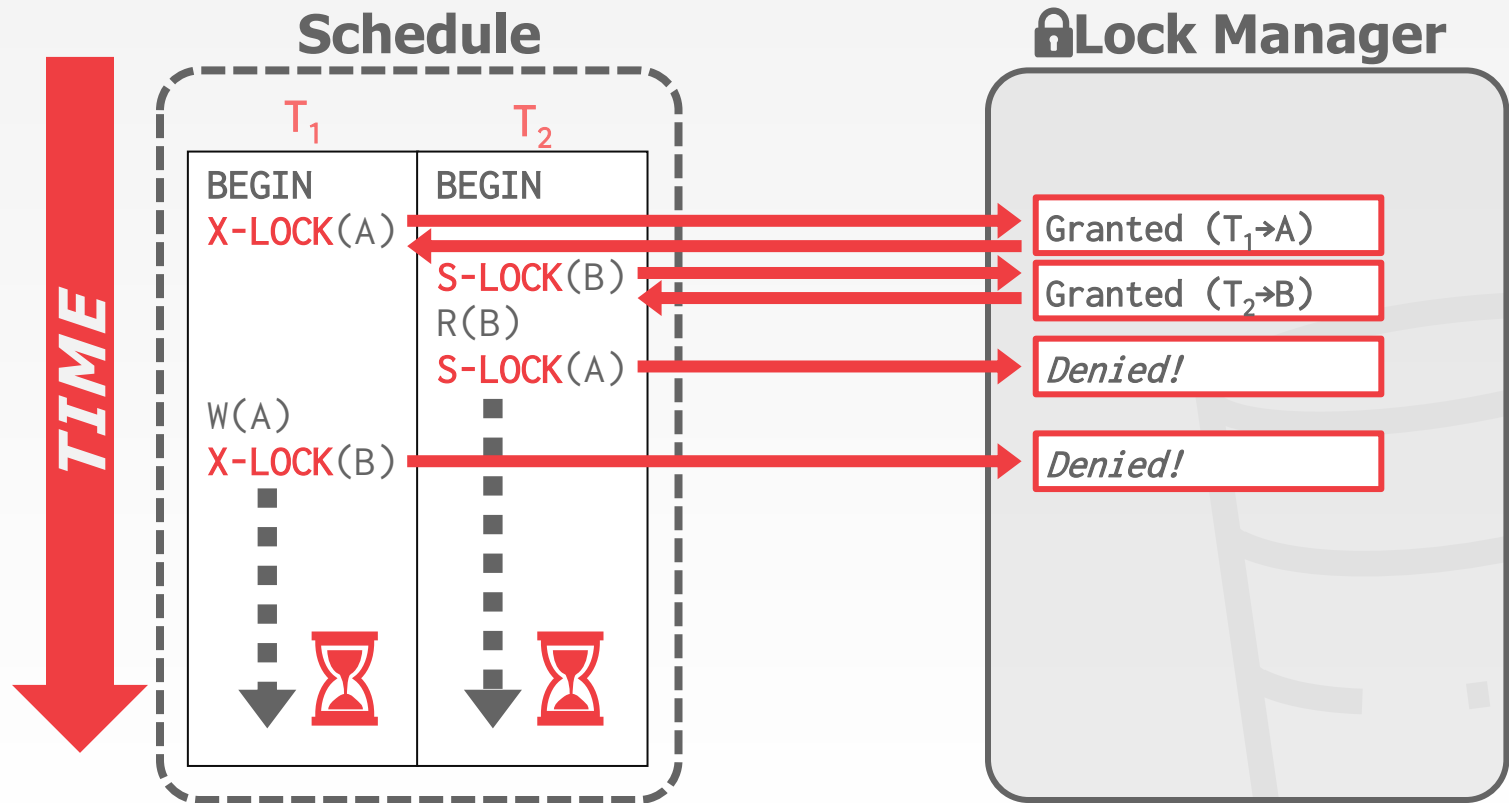
2PL DEADLOCKS



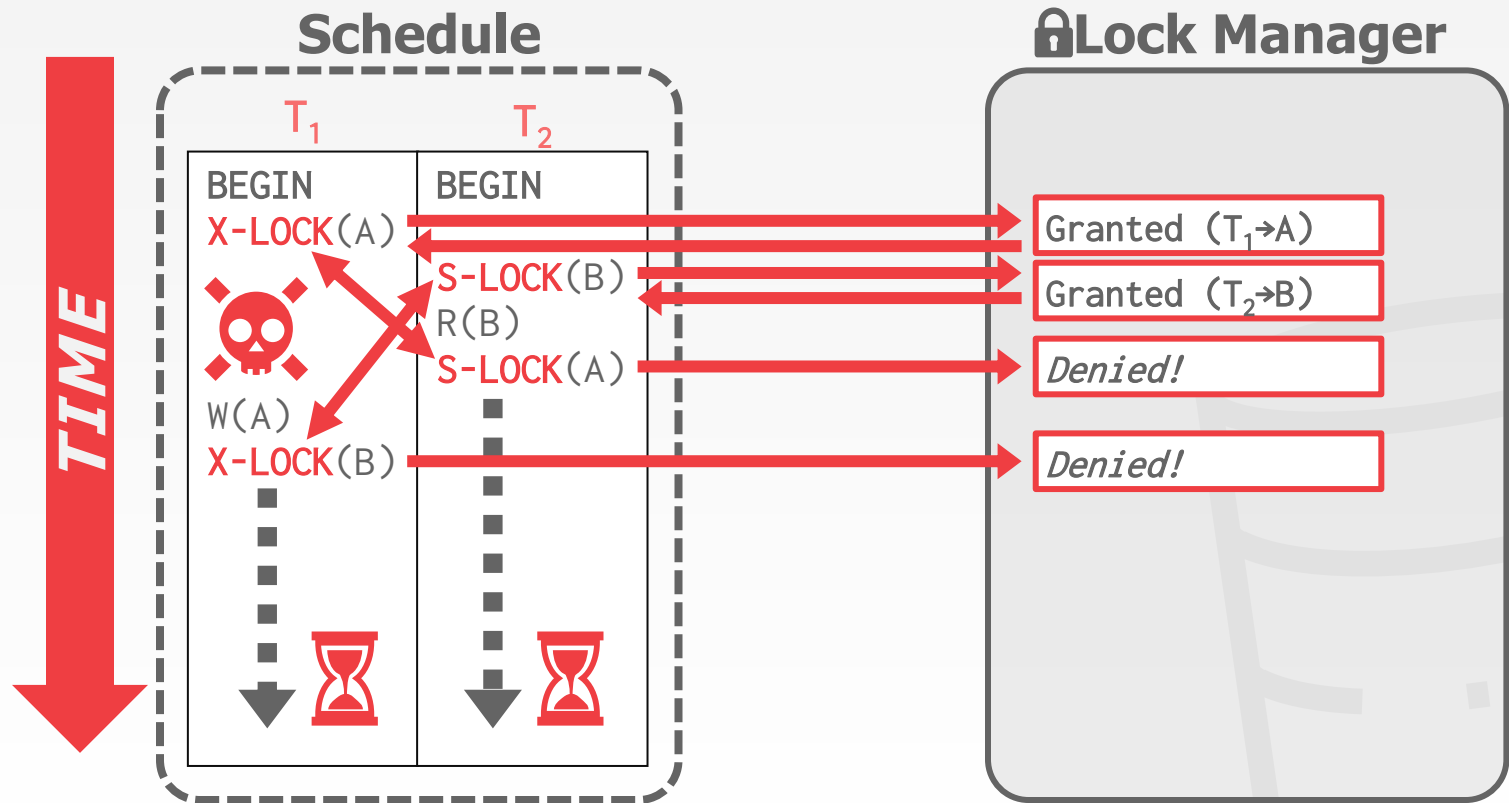
2PL DEADLOCKS



2PL DEADLOCKS



2PL DEADLOCKS



2PL DEADLOCKS

A **deadlock** is a cycle of transactions waiting for locks to be released by each other.

Two ways of dealing with deadlocks:

- **Approach #1: Deadlock Detection**
- **Approach #2: Deadlock Prevention**

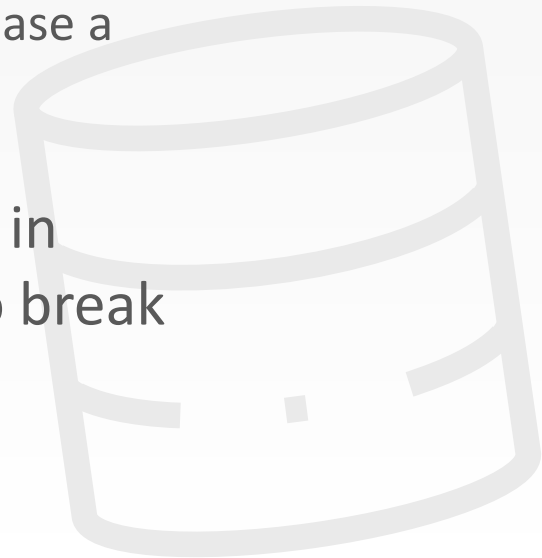


DEADLOCK DETECTION

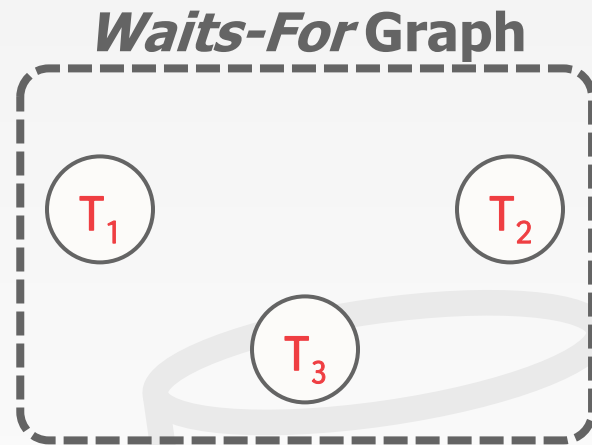
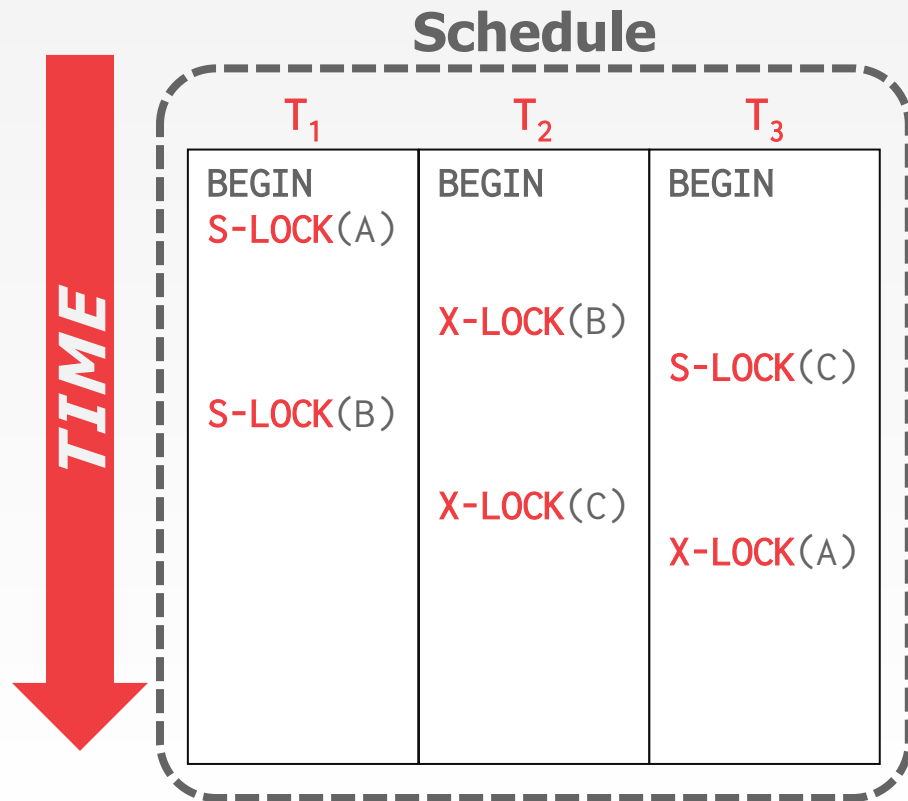
The DBMS creates a **waits-for** graph to keep track of what locks each txn is waiting to acquire:

- Nodes are transactions
- Edge from T_i to T_j if T_i is waiting for T_j to release a lock.

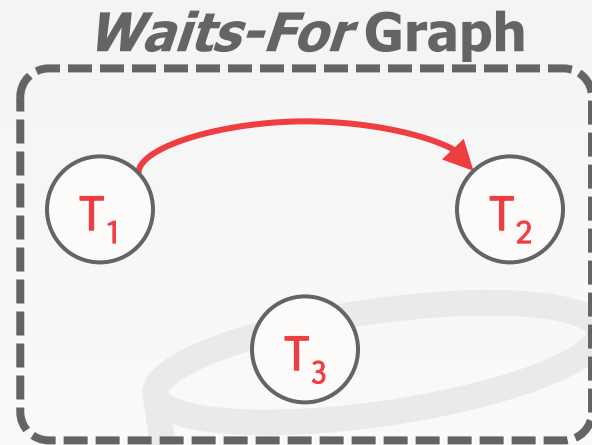
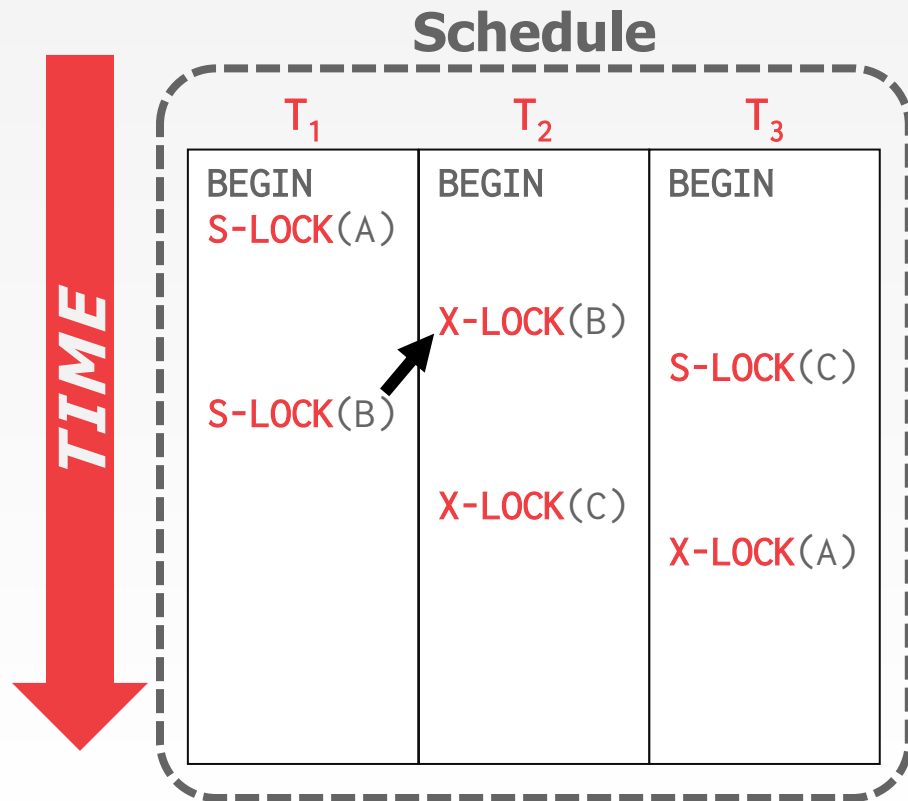
The system periodically checks for cycles in ***waits-for*** graph and then decides how to break it.



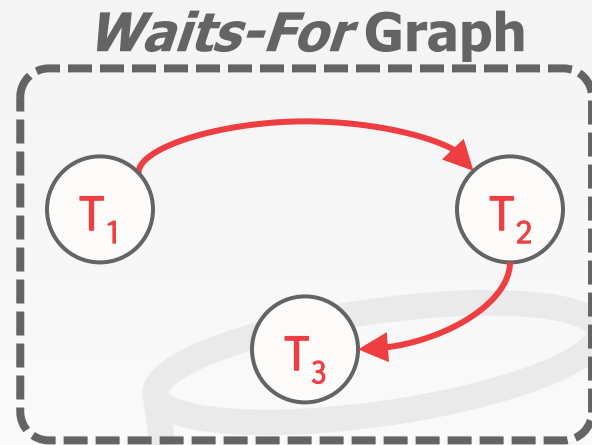
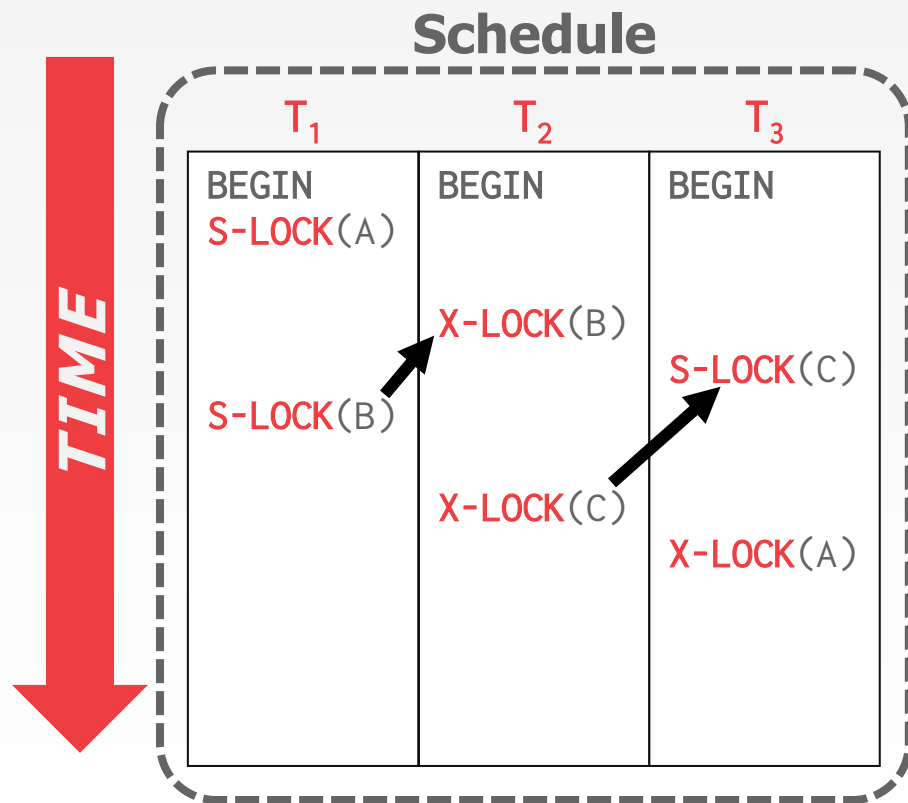
DEADLOCK DETECTION



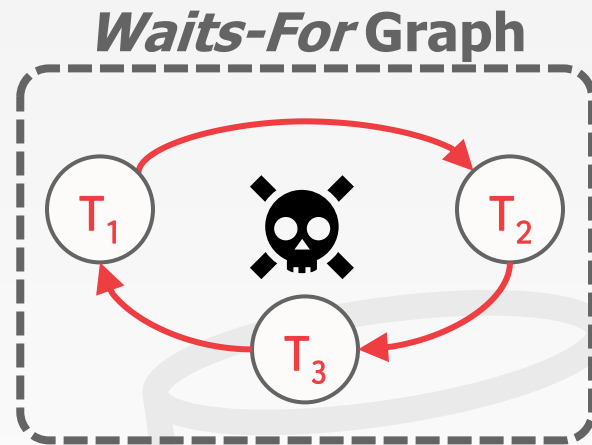
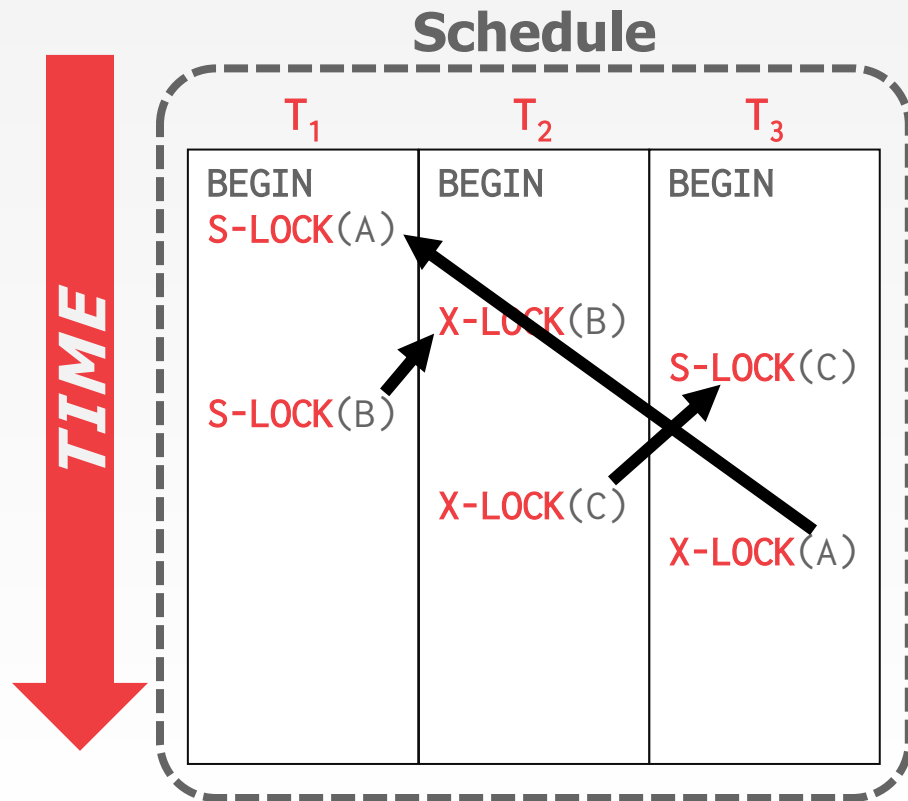
DEADLOCK DETECTION



DEADLOCK DETECTION



DEADLOCK DETECTION

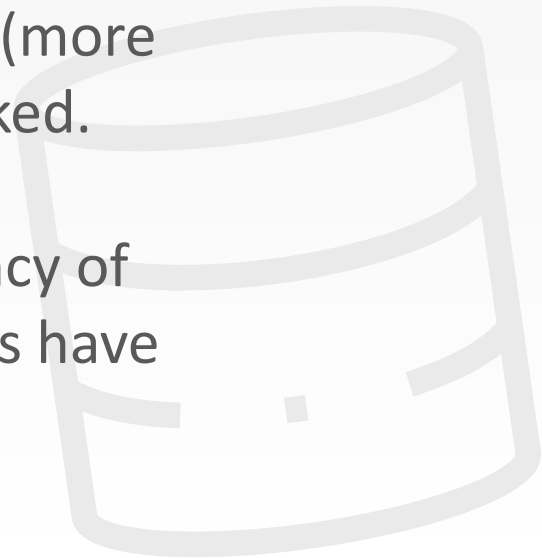


DEADLOCK HANDLING

When the DBMS detects a deadlock, it will select a "victim" txn to rollback to break the cycle.

The victim txn will either restart or abort (more common) depending on how it was invoked.

There is a trade-off between the frequency of checking for deadlocks and how long txns have to wait before deadlocks are broken.



DEADLOCK HANDLING: VICTIM SELECTION

Selecting the proper victim depends on a lot of different variables....



DEADLOCK HANDLING: VICTIM SELECTION

Selecting the proper victim depends on a lot of different variables....

→ By age (lowest timestamp)



DEADLOCK HANDLING: VICTIM SELECTION

Selecting the proper victim depends on a lot of different variables....

- By age (lowest timestamp)
- By progress (least/most queries executed)



DEADLOCK HANDLING: VICTIM SELECTION

Selecting the proper victim depends on a lot of different variables....

- By age (lowest timestamp)
- By progress (least/most queries executed)
- By the # of items already locked



DEADLOCK HANDLING: VICTIM SELECTION

Selecting the proper victim depends on a lot of different variables....

- By age (lowest timestamp)
- By progress (least/most queries executed)
- By the # of items already locked
- By the # of txns that we have to rollback with it

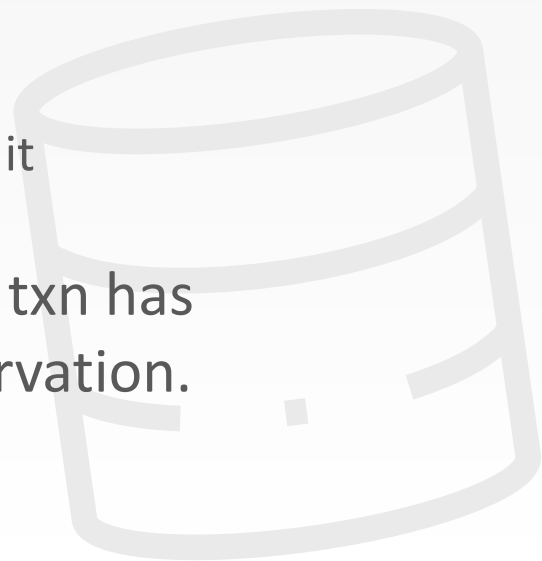


DEADLOCK HANDLING: VICTIM SELECTION

Selecting the proper victim depends on a lot of different variables....

- By age (lowest timestamp)
- By progress (least/most queries executed)
- By the # of items already locked
- By the # of txns that we have to rollback with it

We also should consider the # of times a txn has been restarted in the past to prevent starvation.



DEADLOCK HANDLING: ROLLBACK LENGTH

After selecting a victim txn to abort, the DBMS can also decide on how far to rollback the txn's changes.

Approach #1: Completely

Approach #2: Minimally



DEADLOCK PREVENTION

When a txn tries to acquire a lock that is held by another txn, the DBMS kills one of them to prevent a deadlock.

This approach does not require a ***waits-for*** graph or detection algorithm.



DEADLOCK PREVENTION

Assign priorities based on timestamps:

→ Older Timestamp = Higher Priority (e.g., $T_1 > T_2$)

Wait-Die ("Old Waits for Young")

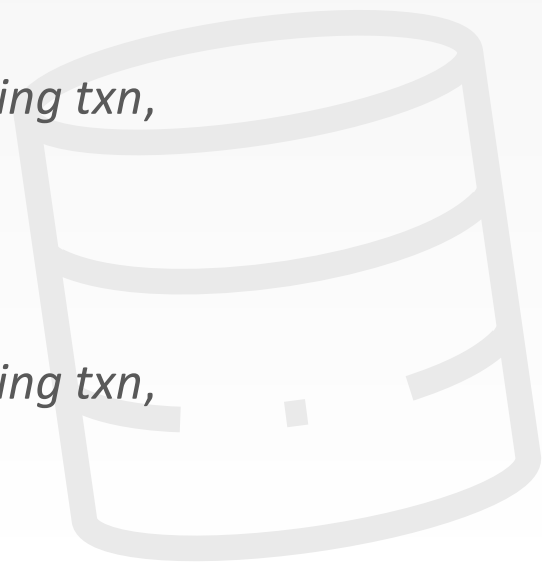
→ If *requesting txn* has higher priority than *holding txn*, then *requesting txn* waits for *holding txn*.

→ Otherwise *requesting txn* aborts.

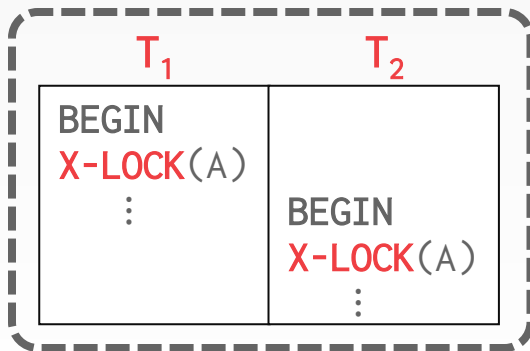
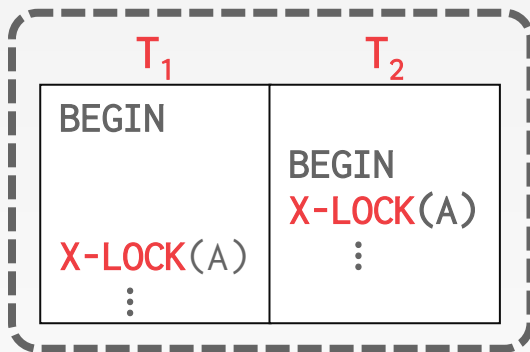
Wound-Wait ("Young Waits for Old")

→ If *requesting txn* has higher priority than *holding txn*, then *holding txn* aborts and releases lock.

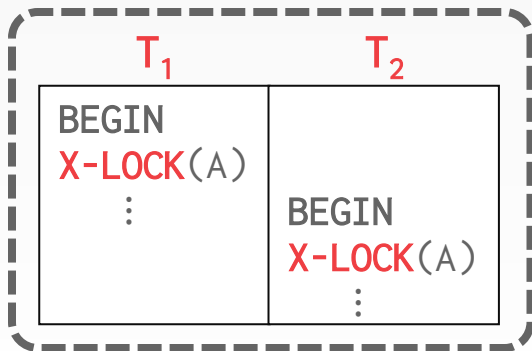
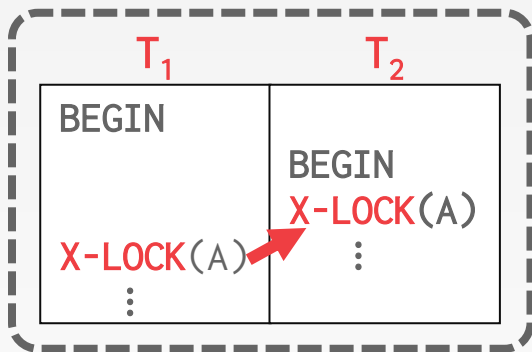
→ Otherwise *requesting txn* waits.



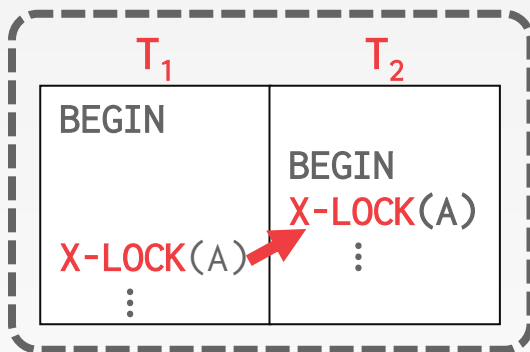
DEADLOCK PREVENTION



DEADLOCK PREVENTION



DEADLOCK PREVENTION

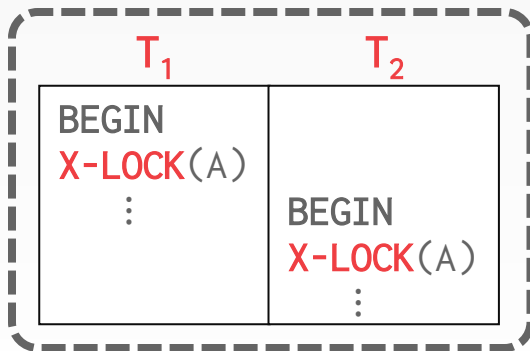


Wait-Die

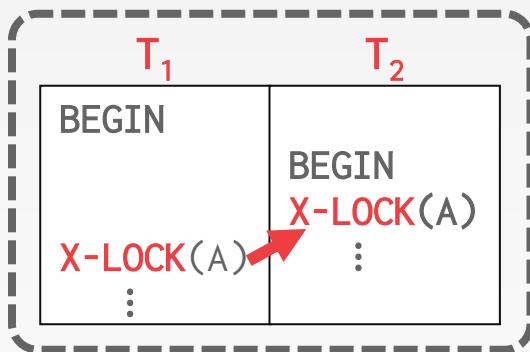
T_1 waits

Wound-Wait

T_2 aborts



DEADLOCK PREVENTION

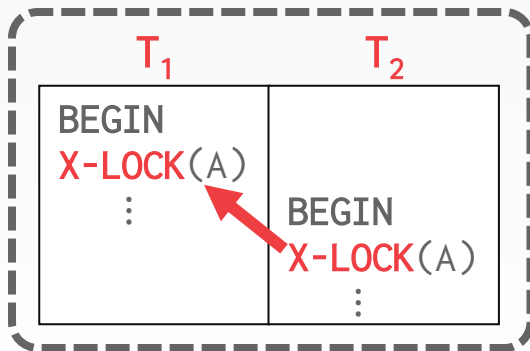


Wait-Die

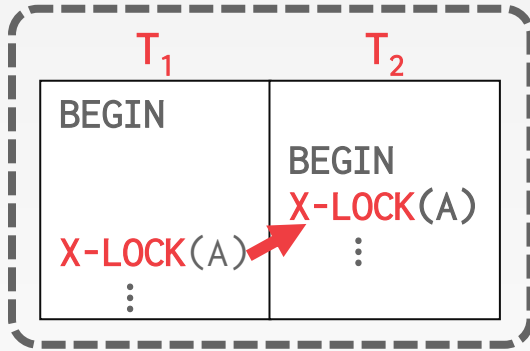
T_1 waits

Wound-Wait

T_2 aborts



DEADLOCK PREVENTION

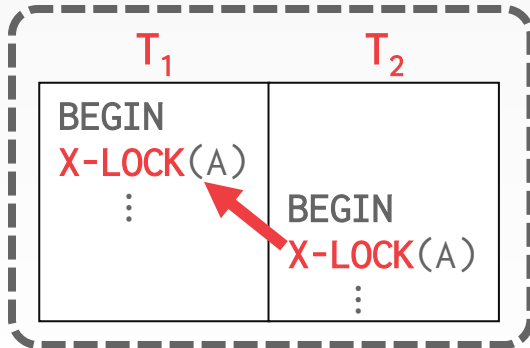


Wait-Die

T_1 waits

Wound-Wait

T_2 aborts



Wait-Die

T_2 aborts

Wound-Wait

T_2 waits



DEADLOCK PREVENTION

Why do these schemes guarantee no deadlocks?

When a txn restarts, what is its (new) priority?



DEADLOCK PREVENTION

Why do these schemes guarantee no deadlocks?

Only one "type" of direction allowed when waiting for a lock.

When a txn restarts, what is its (new) priority?



DEADLOCK PREVENTION

Why do these schemes guarantee no deadlocks?

Only one "type" of direction allowed when waiting for a lock.

When a txn restarts, what is its (new) priority?

Its original timestamp. Why?



WEAKER LEVELS OF ISOLATION

Serializability is useful because it allows programmers to ignore concurrency issues.

But enforcing it may allow too little concurrency and limit performance.

We may want to use a weaker level of consistency to improve scalability.



ISOLATION LEVELS

Controls the extent that a txn is exposed to the actions of other concurrent txns.

Provides for greater concurrency at the cost of exposing txns to uncommitted changes:

- Dirty Reads
- Unrepeatable Reads
- Phantom Reads (Unprotected Inserts/Deletes)



ISOLATION LEVELS

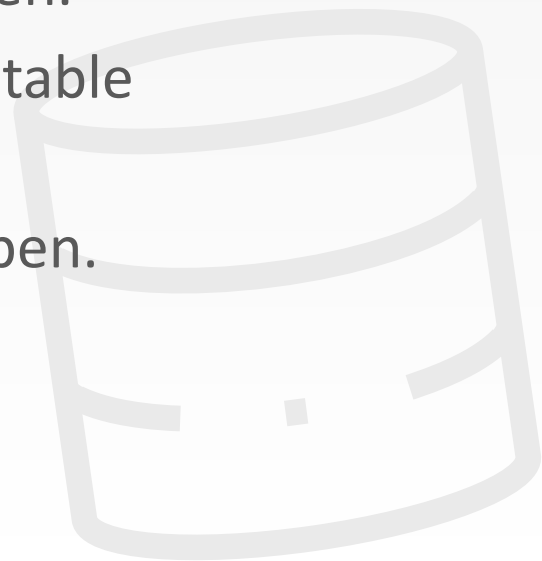


SERIALIZABLE: No phantoms, all reads repeatable, no dirty reads.

REPEATABLE READS: Phantoms may happen.

READ COMMITTED: Phantoms and unrepeatable reads may happen.

READ UNCOMMITTED: All of them may happen.



ISOLATION LEVELS

	<i>Dirty Read</i>	<i>Unrepeatable Read</i>	<i>Phantom</i>
SERIALIZABLE	No	No	No
REPEATABLE READ	No	No	Maybe
READ COMMITTED	No	Maybe	Maybe
READ UNCOMMITTED	Maybe	Maybe	Maybe

SQL-92 ISOLATION LEVELS

You set a txn's isolation level before you execute any queries in that txn.

Not all DBMS support all isolation levels in all execution scenarios

→ Replicated Environments

The default depends on implementation...

```
SET TRANSACTION ISOLATION LEVEL  
<isolation-level>;
```

```
BEGIN TRANSACTION ISOLATION LEVEL  
<isolation-level>;
```



ISOLATION LEVELS (2013)

	<i>Default</i>	<i>Maximum</i>
Action Ingres 10.0/10S	SERIALIZABLE	SERIALIZABLE
Aerospike	READ COMMITTED	READ COMMITTED
Greenplum 4.1	READ COMMITTED	SERIALIZABLE
MySQL 5.6	REPEATABLE READS	SERIALIZABLE
MemSQL 1b	READ COMMITTED	READ COMMITTED
MS SQL Server 2012	READ COMMITTED	SERIALIZABLE
Oracle 11g	READ COMMITTED	SNAPSHOT ISOLATION
Postgres 9.2.2	READ COMMITTED	SERIALIZABLE
SAP HANA	READ COMMITTED	SERIALIZABLE
ScaleDB 1.02	READ COMMITTED	READ COMMITTED
VoltDB	SERIALIZABLE	SERIALIZABLE

Source: [Peter Bailis](#)

ISOLATION LEVELS (2013)

	<i>Default</i>	<i>Maximum</i>
Action Ingres 10.0/10S	SERIALIZABLE	SERIALIZABLE
Aerospike	READ COMMITTED	READ COMMITTED
Greenplum 4.1	READ COMMITTED	SERIALIZABLE
MySQL 5.6	REPEATABLE READS	SERIALIZABLE
MemSQL 1b	READ COMMITTED	READ COMMITTED
MS SQL Server 2012	READ COMMITTED	SERIALIZABLE
Oracle 11g	READ COMMITTED	SNAPSHOT ISOLATION
Postgres 9.2.2	READ COMMITTED	SERIALIZABLE
SAP HANA	READ COMMITTED	SERIALIZABLE
ScaleDB 1.02	READ COMMITTED	READ COMMITTED
VoltDB	SERIALIZABLE	SERIALIZABLE

Source: [Peter Bailis](#)

ISOLATION LEVELS (2013)

	<i>Default</i>	<i>Maximum</i>
Action Ingres 10.0/10S	SERIALIZABLE	SERIALIZABLE
Aerospike	READ COMMITTED	READ COMMITTED
Greenplum 4.1	READ COMMITTED	SERIALIZABLE
MySQL 5.6	REPEATABLE READS	SERIALIZABLE
MemSQL 1b	READ COMMITTED	READ COMMITTED
MS SQL Server 2012	READ COMMITTED	SERIALIZABLE
Oracle 11g	READ COMMITTED	SNAPSHOT ISOLATION
Postgres 9.2.2	READ COMMITTED	SERIALIZABLE
SAP HANA	READ COMMITTED	SERIALIZABLE
ScaleDB 1.02	READ COMMITTED	READ COMMITTED
VoltDB	SERIALIZABLE	SERIALIZABLE

Source: [Peter Bailis](#)

SQL-92 ACCESS MODES

You can provide hints to the DBMS about whether a txn will modify the database during its lifetime.

Only two possible modes:

→ **READ WRITE** (Default)

→ **READ ONLY**

Not all DBMSs will optimize execution if you set a txn to in **READ ONLY** mode.

```
SET TRANSACTION <access-mode>;
```

```
BEGIN TRANSACTION <access-mode>;
```



CONCLUSION

2PL is used in almost DBMS.

Automatically generates correct interleaving:

- Locks + protocol (2PL, SS2PL ...)
- Deadlock detection + handling
- Deadlock prevention

Default DBMS isolation levels are usually weaker than serializable



NEXT CLASS

Logging

